

Do not distribute course material

You may not and may not allow others to reproduce or distribute lecture notes and course materials publicly whether or not a fee is charged.

Merge sort is used in
data bases!
It is a key component in
Hadoop, a distributed file
system.

Merge Sort continued

Merge Sort

43	37	20	25	4	72	15	19
----	----	----	----	---	----	----	----

DIVIDE IN HALF

43	37	20	25	4	72	15	19
----	----	----	----	---	----	----	----

CONQUER: SORT EACH HALF

20	25	37	43	4	15	19	72
↑				↑			

left right

COMBINE: MERGE

4	15	19	20	25	37	43	72
---	----	----	----	----	----	----	----

Pseudocode

MERGE-SORT(A, p, r)

if $p < r$ ← check for base case
 $q = \lfloor(p + r)/2\rfloor$ ← divide
 MERGE-SORT(A, p, q) ← conquer
 MERGE-SORT($A, q+1, r$) ← conquer
 MERGE(A, p, q, r) ← combine

A

20	25	20	29	42	43	49	79
----	----	----	----	----	----	----	----

MERGE-SORT($A, 1, 1$)

MERGE-SORT($A, 1, 2$)

MERGE-SORT($A, 1, 4$)

MERGE-SORT($A, 1, 8$)

MERGE(A, p, q, r)

$$n_1 = q - p + 1$$

$$n_2 = r - q$$

let $L[1..n_1 + 1]$ and $R[1..n_2 + 1]$ be new arrays

for $i = 1$ **to** n_1

$$L[i] = A[p + i - 1]$$

for $j = 1$ **to** n_2

$$R[j] = A[q + j]$$

$$L[n_1 + 1] = \infty$$

$$R[n_2 + 1] = \infty$$

$$i = 1$$

$$j = 1$$

for $k = p$ **to** r

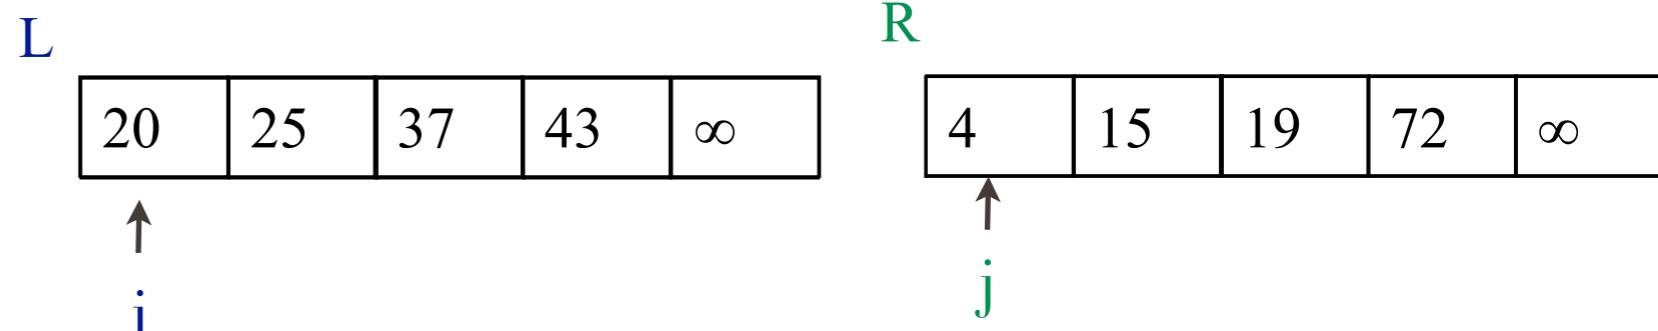
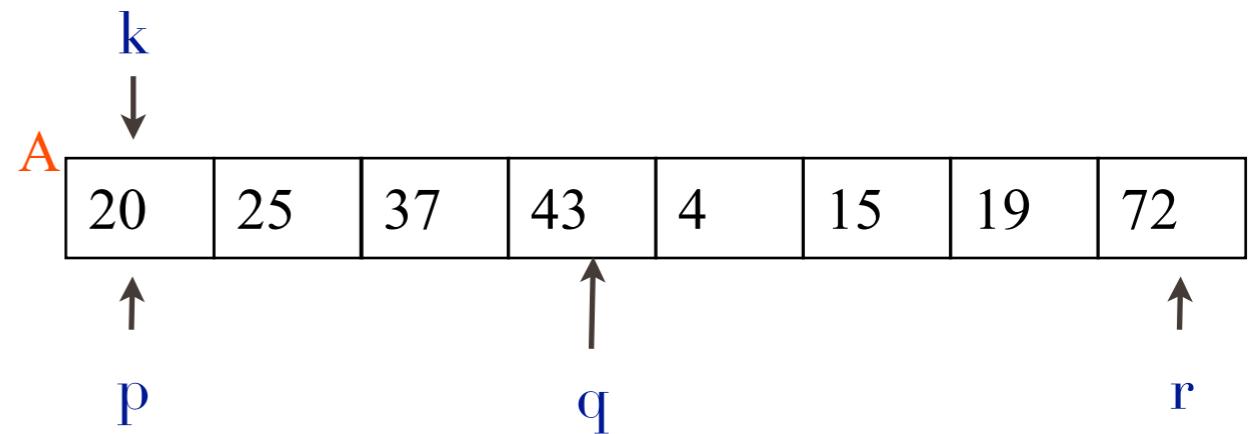
if $L[i] \leq R[j]$

$$A[k] = L[i]$$

$$i = i + 1$$

else $A[k] = R[j]$

$$j = j + 1$$



Every time I go through the loop I perform a constant amount of work

See page 31 and 33 in the textbook for a proof that the merge function works correctly

Pseudocode

```
MERGE-SORT(A, p, r)           T(n) =  
    if p < r  
        q = ⌊(p + r)/2⌋          }   Θ(1)  
        MERGE-SORT(A, p, q)      }   2 · T(n / 2)  
        MERGE-SORT(A, q+1, r)  }  
        MERGE(A, p, q, r)       Θ(n)
```

Running time: $T(n) = ?$

$$T(1) = \Theta(1)$$

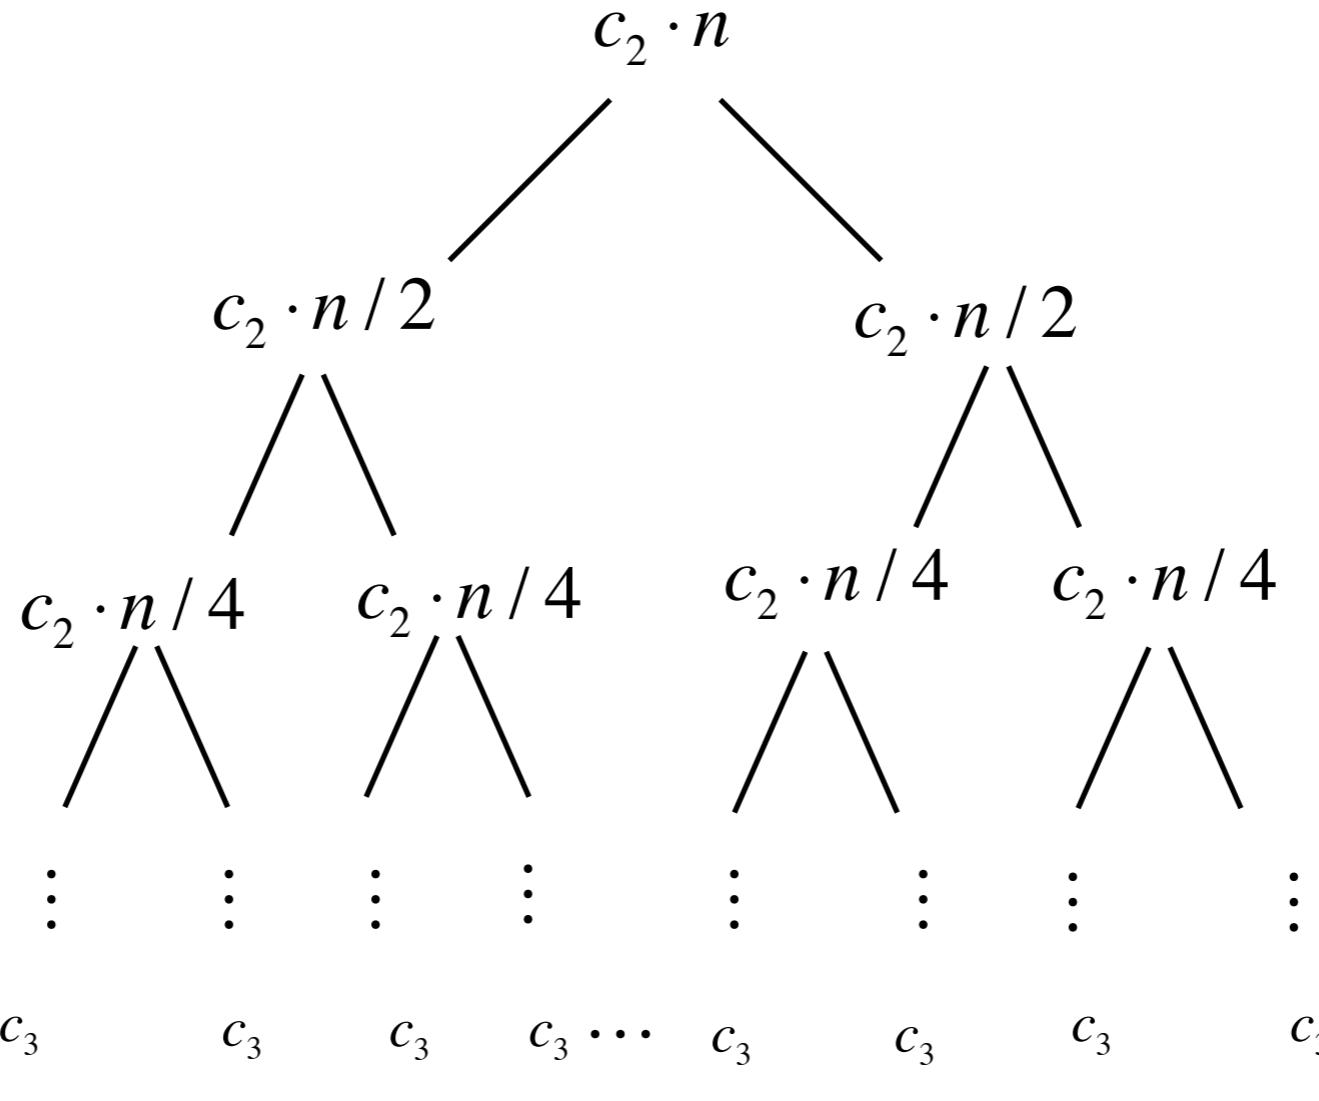
$$T(n) = \Theta(1) + 2 \cdot T(n / 2) + \Theta(n)$$

$$T(n) = \begin{cases} \Theta(1) & \text{if } n = 1, \\ 2T(n / 2) + \Theta(n) & \text{otherwise} \end{cases}$$

Here for simplicity, we assume $n = 2^n$

$$T(n) = \begin{cases} c_3 & \text{if } n = 1, \\ 2T(n/2) + c_2 \cdot n & \text{otherwise} \end{cases}$$

$1 + \log n$ levels



# per level	level cost
2^0	$1 \cdot c_2 \cdot n$
2^1	$2^1 c_2 \cdot \frac{n}{2}$
2^2	$2^2 c_2 \cdot \frac{n}{2^2}$
:	:
$2^{\log_2 n} = n$	$n \cdot c_3 \cdot \frac{n}{n}$

Total $c_3 \cdot n + c_2 \cdot n \cdot \log n$

therefore $T(n)$ is $\Theta(n \log n)$

Pseudocode

MERGE-SORT(A, p, r) $T(n) =$

if $p < r$ } $\Theta(1)$
 $q = \lfloor (p + r)/2 \rfloor$ }
 MERGE-SORT(A, p, q) } $2 \cdot T(n/2)$
 MERGE-SORT(A, q+1, r) }
 MERGE(A, p, q, r) $\Theta(n)$

Running time: $T(n)$ is $\Theta(n \log n)$

Hot shot Programmer/amazing machine vs
a good algorithm/average machine/bad compiler

insertion sort merge sort

Computer A

Best computer

10 billion instructions per second

Hot shot programmer

$2n^2$ instructions to sort n numbers

n

10

10 million

100 million

$$\frac{2 * (10)^2 \text{ instructions}}{10^{10} \text{ instructions/second}}$$

$$\frac{2 * (10^7)^2 \text{ instructions}}{10^{10} \text{ instructions/second}}$$

$$\frac{2 * (10^9)^2 \text{ instructions}}{10^{10} \text{ instructions/second}}$$

Computer B

ok computer

10 million instructions per second

Average programmer

(inefficient compiler)

50 $n \log(n)$ instructions to sort n numbers

$$\frac{50 * 10 \lg(10) \text{ instructions}}{10^7 \text{ instructions/second}}$$

$$\frac{50 * 10^7 \lg(10^7) \text{ instructions}}{10^7 \text{ instructions/second}}$$

$$\frac{50 * 10^9 \lg(10^9) \text{ instructions}}{10^7 \text{ instructions/second}}$$

insertion sort merge sort

Computer A

Best computer

10 billion instructions per second

Hot shot programmer

$2n^2$ instructions to sort n numbers

Computer B

ok computer

10 million instructions per second

Average programmer

(inefficient compiler)

50 $n \log(n)$ instructions to sort n numbers

n

10

20 nanoseconds

166.1 microseconds

10 million

5 hours 33 min. 20 seconds 19 min. 22.67 seconds

100 million

6 years 4 months 3.1 days 1 day 17 hours 31 min. 22.76 seconds

Running Times

n	insertion $O(n^2)$	merge $O(n \log(n))$
512	0.0018	0.000166
1024	0.00747	0.000356
2048	0.030801	0.000765
4096	0.120905	0.001644
8192	0.474183	0.003504
16384	1.87247	0.007584
32768	7.56883	0.015684
65536	29.9251	0.033017

Asymptotic Notation Examples

- | | |
|------------------------------------|----------------------------|
| 1. $\log_{200} n$ | $\in \Theta(\log_2 n)$ |
| 2. $\log n!$ | $\in \Theta(n \log n)$ |
| 3. $n! + C^n$ | $\in \Theta(n!)$ |
| 4. $n! + n^n$ | $\in \Theta(n^n)$ |
| 5. $\log^{200} n = (\log n)^{200}$ | $\in \Theta(\log^{200} n)$ |

Useful Facts:

Change of basis

$$\log_a b = \log_a c^{\log_c b} = \log_c b \log_a c$$

Stirling's approximation

$$n! = \sqrt{2\pi n} \left(\frac{n}{e} \right)^n \left(1 + O\left(\frac{1}{n}\right) \right)$$

Some Important Big-Oh's

• $O(1)$	constant	assignment statement invoking a function
• $O(\log n)$	logarithmic	binary search inserting into a red-black tree
• $O(n)$	linear	searching in an unordered list inserting a new item into a sorted vector
• $O(n \log n)$	$n \log n$	mergesort
• $O(n^2)$	quadratic	insertion sort
• $O(n^3)$	cubic	Brute force maximum subsequence sum algorithm
• $O(2^n)$	exponential	Finding all subsets of n items

Elementary Data Structures

It is all about the data!

- Basic Abstract Data Types (ADTs)
- Implementing a Priority Queue
 - Binary Heap
- Heap Sort

Data Structure

“A **data structure** is a way to store and organize data in order to facilitate access and modifications.
No single data structure works well for all purposes, ...”

ADT (Abstract Data Type)

The Specification

A mathematical model of the data
methods used to modify and access the data
~~how the data organized in memory~~
~~what algorithms implement the methods~~

ADT

When designing an algorithm we use the behavior of the ADT *without thinking* of how it is implemented.

We have different implementations of the ADT, so we can *optimize* for performance

A large yellow five-pointed star shape is positioned in the upper right corner of the slide. It has a white center and a thin black outline.

We chose
the data structure that works
best for the
algorithm.

Do we just need to insert, delete
and test membership?

Do we need to insert and find the
minimum item?

ADT - Abstract Data Types

- list, stack, queue, priority queue, set, dictionary

Common data structures

- array, linked list, hash table, tree (binary heap, binary search tree, red-black tree, 2-3 tree, b-tree)

“Most programmers have seen them, and most good programmers realize they’ve written at least one. They are huge, messy, ugly program that should have been short, clean, beautiful programs.”

“Why do programmers write big programs when small ones will do?
... arrays are typically used as fixed tables that are initialized at the beginning of a program and never altered.”

We often use these when we design an algorithm



Data Structure Needed For:

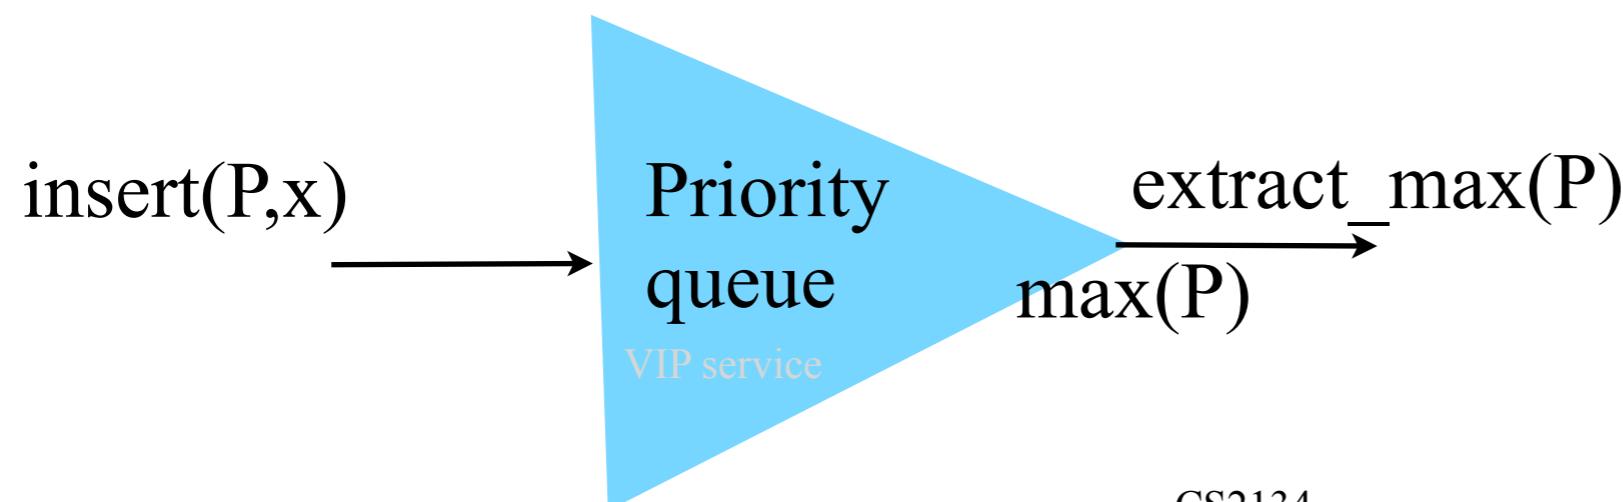
<i>free list</i> used to keep track of free memory in the heap (free store)	Determining if any of a company's customers is on a list of hacked SSN
Order fulfillment system for a company	Finding a name in a telephone book, and all the people with the same last name
Buffering video as it is streamed	Storing the address of recently visited web sites
Scheduling airplanes landing on Newark's runway	Creating a continuous limit order book (buyers post bids for stocks, and seller post offers to sell)
Online ticketing requests	

ADT's

stack, queue, priority queue

Retrieve items as a function of when it arrived, or its priority

$\text{push}(S, x)$ $\text{pop}(S)$

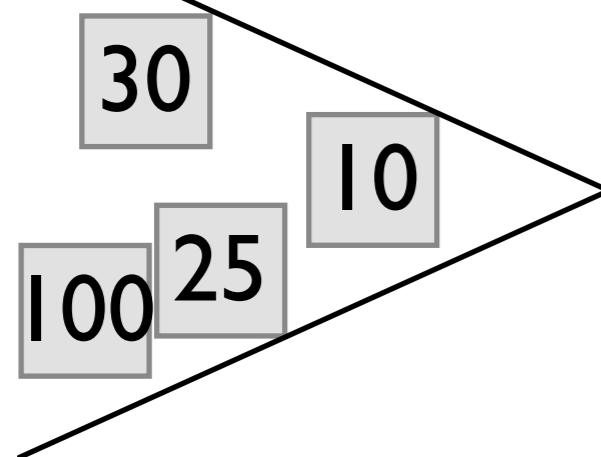


Which ADT data structure should you use?

ADT's

stack, queue, priority queue

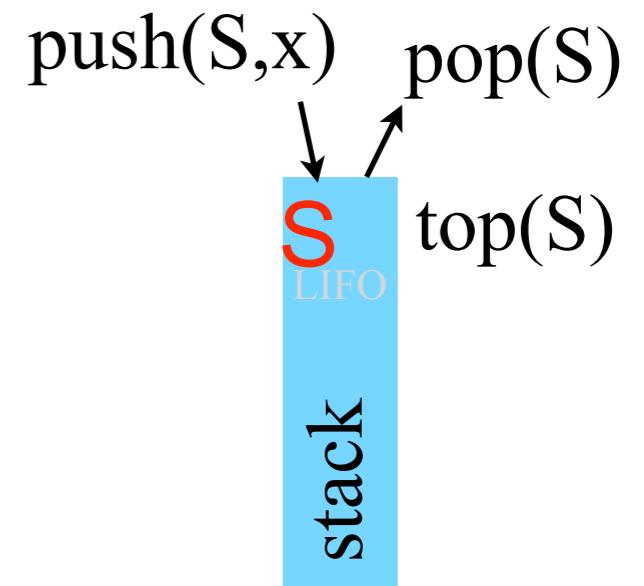
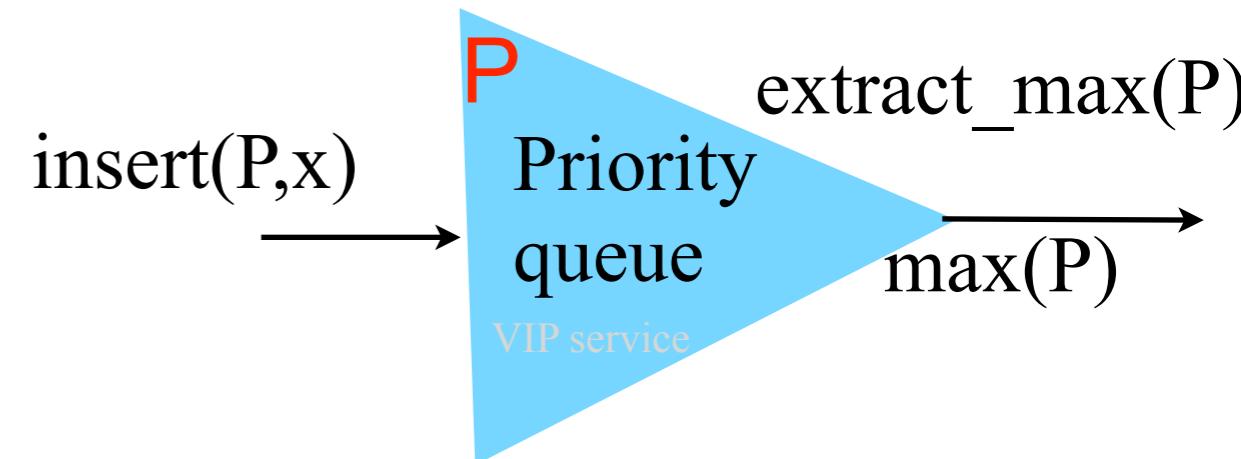
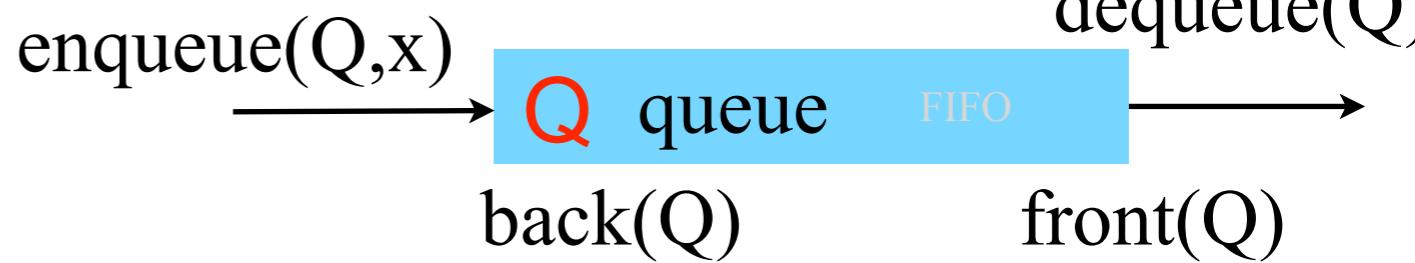
Printer Jobs



Function Calls



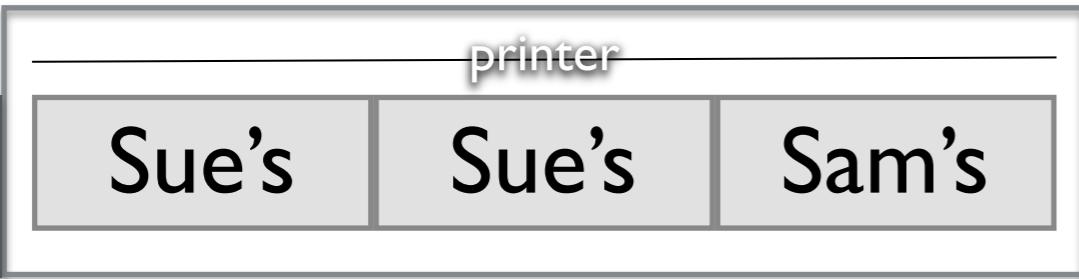
customersInLine



computer 1

printer queue

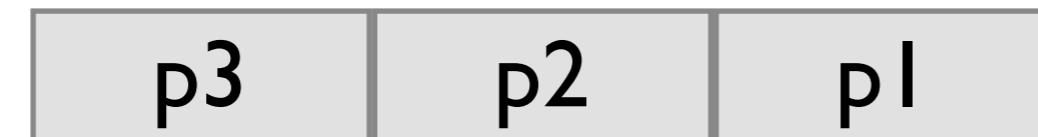
computer 2



video buffering

stored video file

video player



Back Button for Web Pages

nyu.ed

yahoo.com

google.com

Which ADT data structure should you use?

ADT's

stack, queue, priority queue

push(x) pop()



Running Time for this implementation?

Push(S,x)

Pop(S)

Stack-Empty(S)

When poll is active, respond at **PollEv.com/lindamsellie089**

Text **LINDAMSELLIE089** to **22333** once to join

How long does it take to enter a new item into a stack in the worst case (assuming there is room in the array implementation)?

$O(1)$

$O(\log n)$

$O(n^{1/2})$

$O(n)$

None of the above

Total Results: 0

ADT's

stack, queue, priority queue

Implementation using an array



Instance attributes: `Q.head`, `Q.tail`, `Q.length`

Running Time for this implementation?

`Enqueue(Q,x)`

`Dequeue(Q)`

`Queue-Empty(Q)`

Q.head is the index of the most oldest item in the queue
Q.tail is the index just after the most recently inserted item
Q[Q.head .. Q.tail-1] is the queue where we “wrap around”
Q.length is the size of the array

When poll is active, respond at **PollEv.com/lindamsellie089**

Text **LINDAMSELLIE089** to **22333** once to join

In the worst case, how long does it take to remove an item from a queue?

$O(1)$

$O(\log n)$

$O(n)$

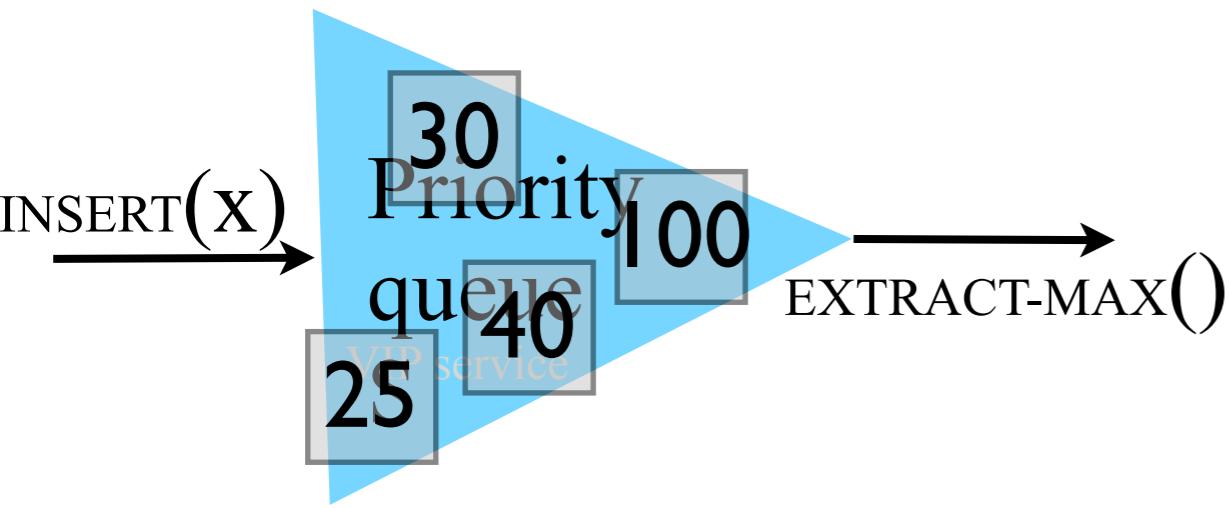
$O(n \log n)$

None of the above

Total Results: 0

ADT Priority Queue

Stores data that is associated with a **key**
(key from a linearly ordered set)



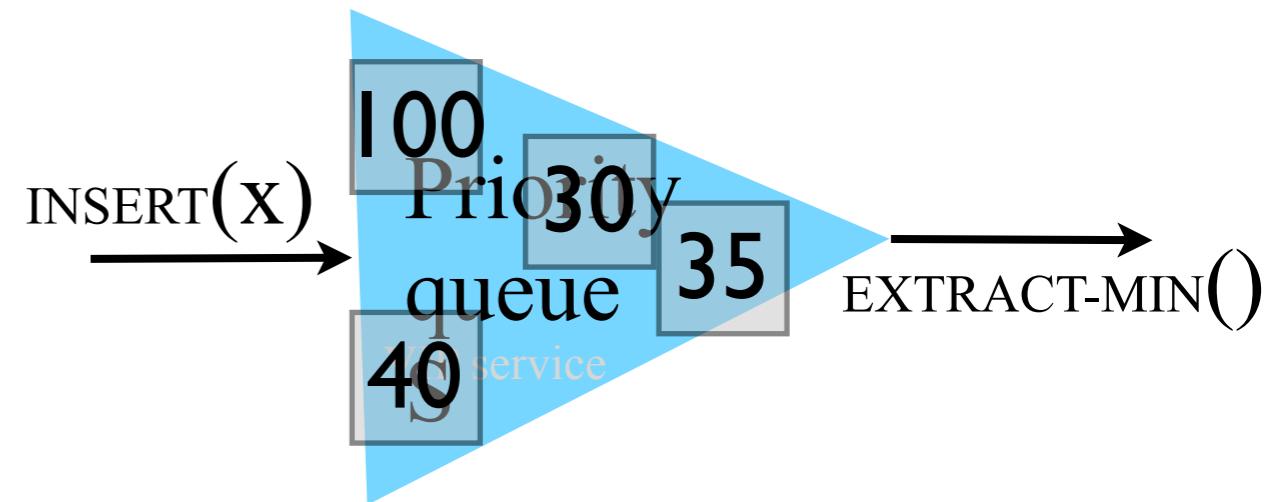
max-priority queue

INSERT(S,x)

MAXIMUM(S)

EXTRACT-MAX(S)

INCREASE-KEY(S,x,k)



min-priority queue

INSERT(S,x)

MINIMUM(S)

EXTRACT-MIN(S)

DECREASE-KEY(S,x,k)

Applications

- Operating systems - process jobs by priority
- Graph algorithms - subroutine in a Dijkstra's shortest path algorithm
- Event simulation - find next event to happen
- Compression - Huffman encoding
- Heapsort

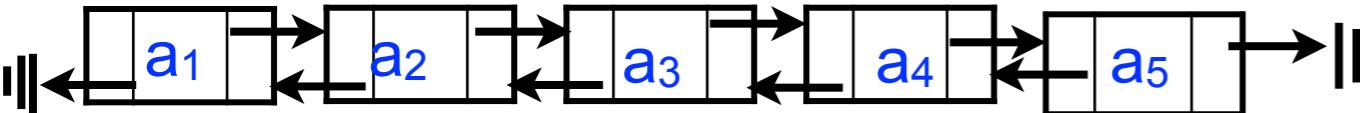
Pair share (write down the question)

How should we implement the priority queue?



- sorted list (array)?

- unsorted list(array)?



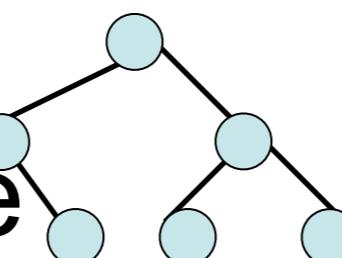
- sorted list (linked list)?

- unsorted list(linked list)?

- queue?

- stack?

- binary search tree



- balance binary search

INSERT(S,x)

EXTRACT-MIN(S)

Priority Queues are often implemented with a **binary heap**

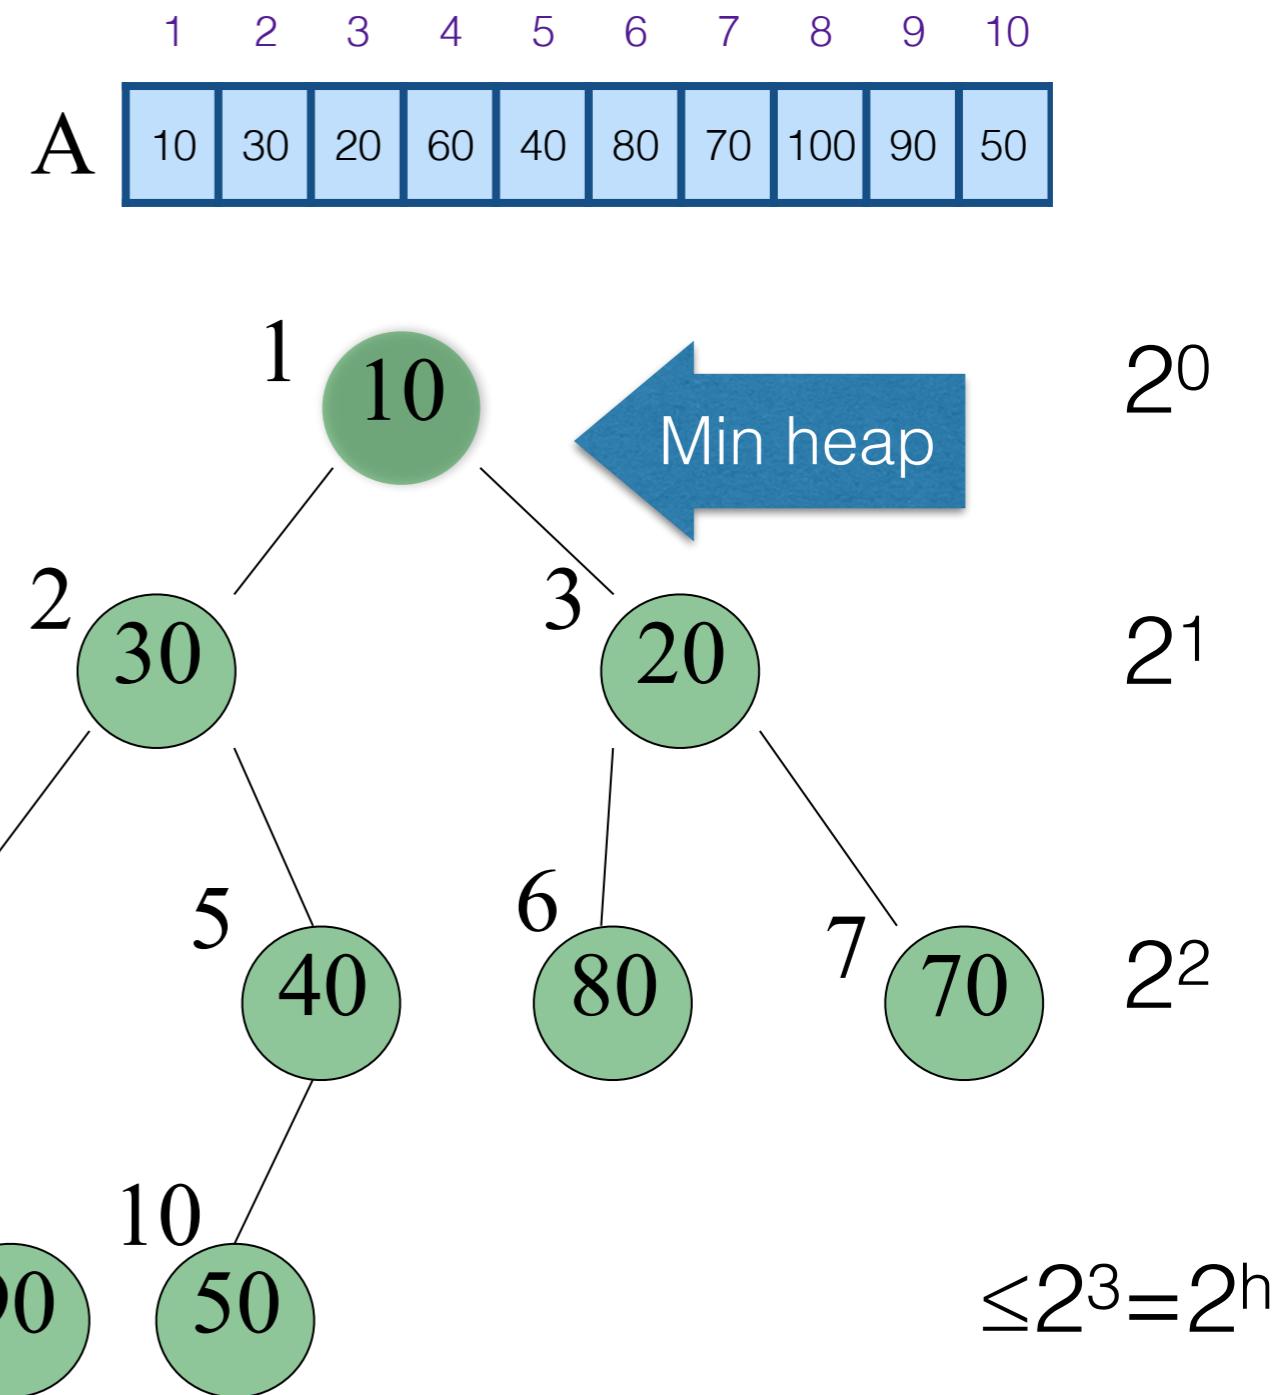
Each element in the array represents a tree node. How?

Elements stored in
A[1..heap-size]

Binary Heap

Attributes:
length
heap-size

- array structure
- visualized as a *nearly complete tree*
- root of tree is $A[1]$
- $\text{PARENT}(i) = i/2$
- $\text{LEFT}(i) = 2i$
- $\text{RIGHT}(i) = 2i + 1$
- computing is fast with binary representation



What is the maximum number of nodes in a tree of height h?

2^h

2^h

$2^h + 1$

2^{h-1}

$2^{h-1} + 1$

$2^{h+1} - 1$

None of the above

Theorem:

Max heap

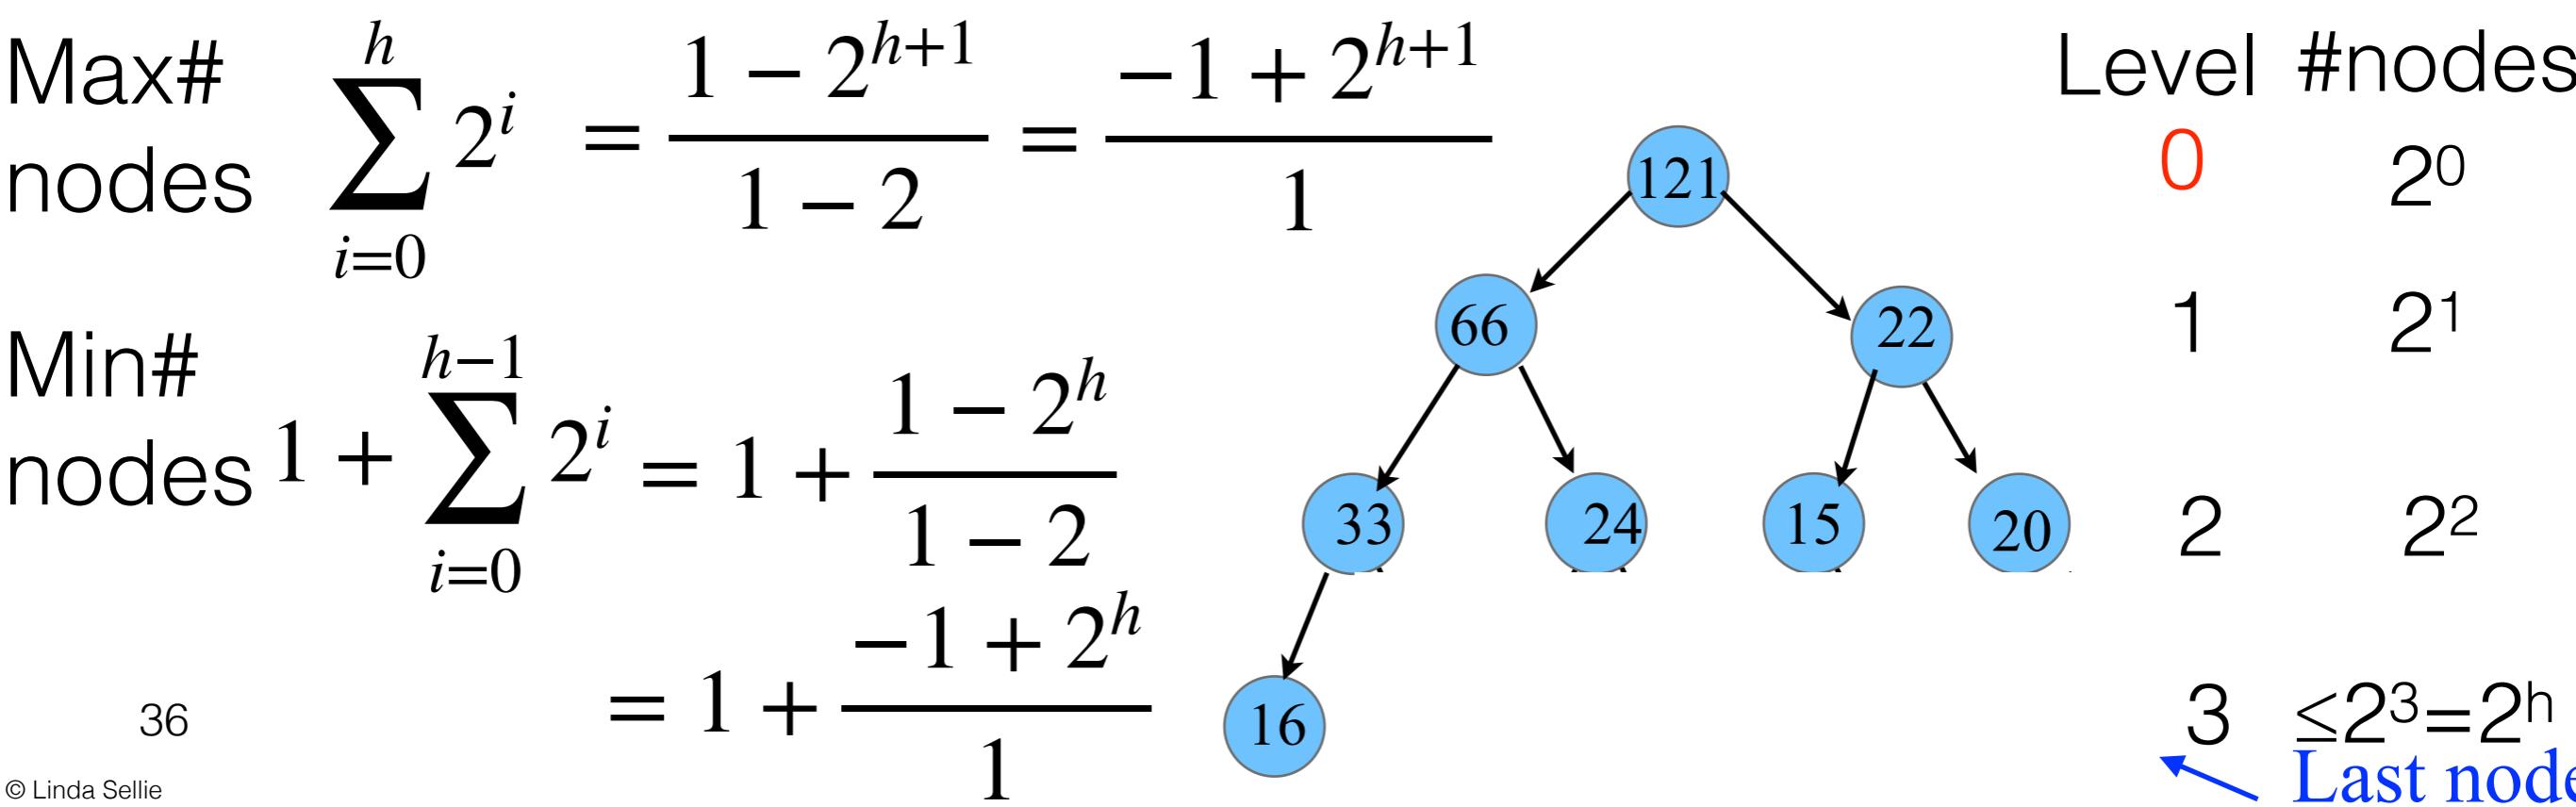
A nearly complete binary tree of n nodes has height $\Theta(\log n)$.

Proof:

The number of nodes in a complete binary tree where all levels have the maximum number of nodes is $2^{h+1}-1$ where h is the height of the tree.

Therefore the number of nodes of a nearly complete tree of height h is $2^h \leq n \leq 2^{h+1}-1$

$$h \leq \log_2 n < h + 1$$



Binary Heap Data Structure

- **Nearly Complete binary tree:** binary tree in which all levels are “full”, except for possibly bottom level which is filled in from the left
- Where h is height and n is size is $h = \Theta(\log n)$
- A **min binary heap** (min-heap for this lecture) is a complete binary tree in which every node satisfies the heap property:
 - $\text{PARENT}(A[i]) \leq A[i]$
- A **max binary heap** (max-heap) is a complete binary tree in which every node satisfies the heap property:
 - $\text{PARENT}(A[i]) \geq A[i]$
- Observe
 - By induction and transitivity of \leq , in a min heap, the smallest node is at the root. Similarly for a max heap
 - heap is NOT Binary Search Tree

“Heap may refer to:

Computer science

- Heap (data structure), a data structure commonly used to implement a priority queue
- Heap (or *free store*), an area of memory used for dynamic memory allocation

Mathematics

- Heap (mathematics), a generalization of a group”

From <http://en.wikipedia.org/wiki/Heap>

Min-heap operations

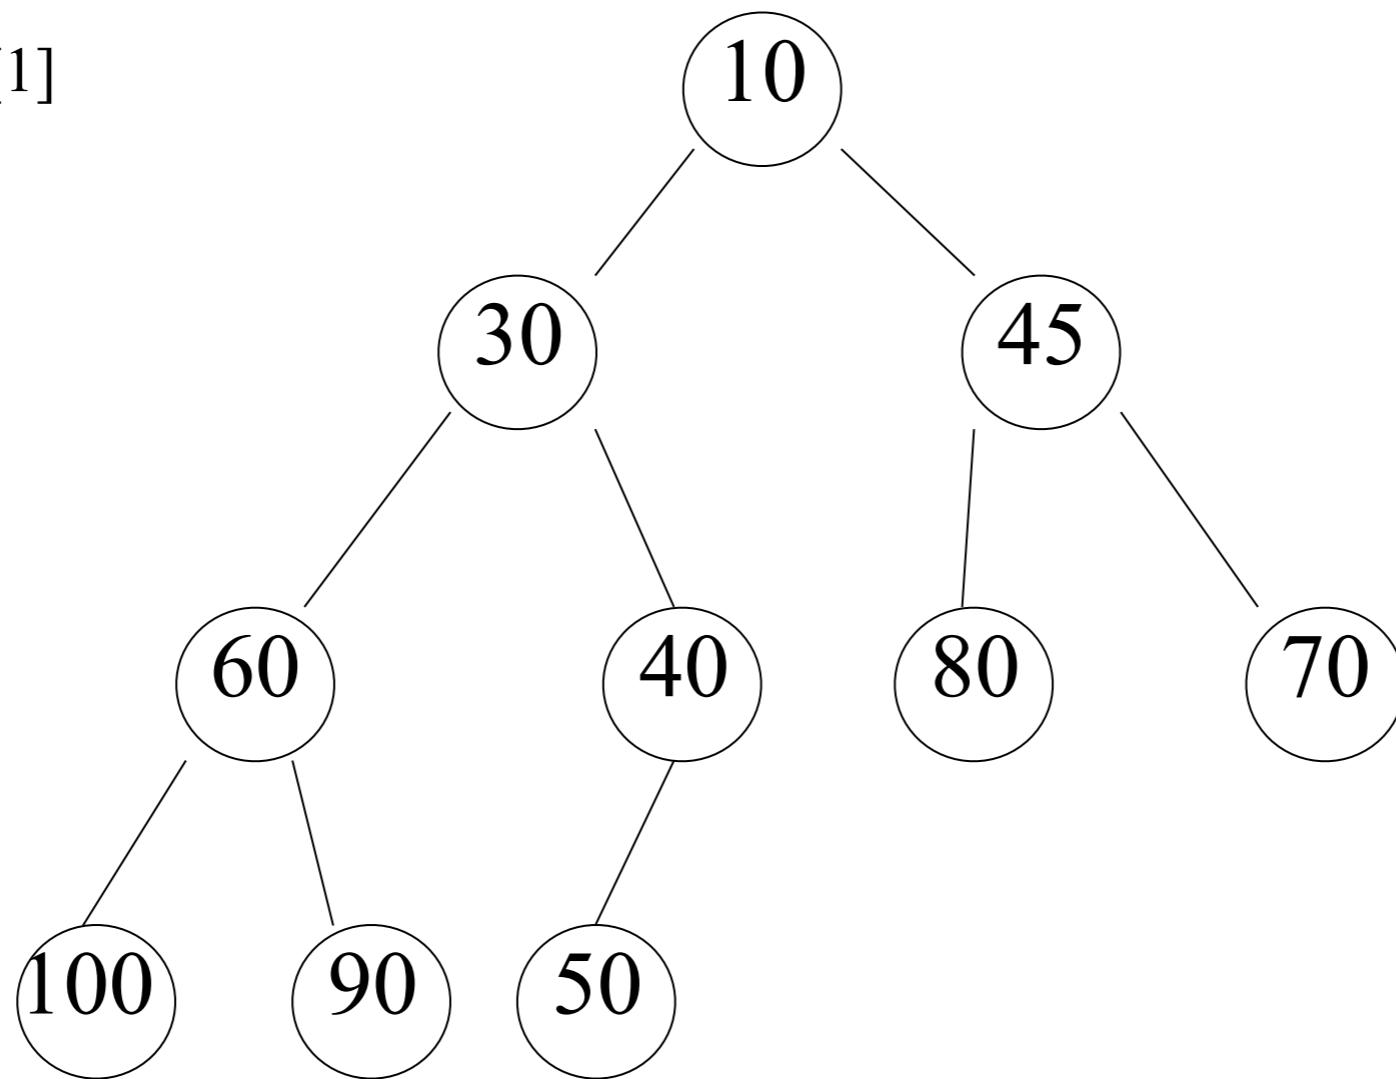
- HEAP-MINIMUM(A)
- HEAP-EXTRACT-MIN(A)
- MIN-HEAP-INSERT(A,key)
- HEAP-DECREASE-KEY(A, i, key)
- BUILD-MIN-HEAP(A) // build a heap from an unsorted array
- MIN-HEAPIFY(A,i)

A max heap works the same as a min heap:
exchange min for max
and when comparing keys
exchange $<$ for $>$
and \leq for \geq

Finding the Minimum item in a Min Heap

HEAP-MINIMUM(A)

return A[1]



10	30	45	60	40	80	70	100	90	50
----	----	----	----	----	----	----	-----	----	----

1	2	3	4	5	6	7	8	9	10
---	---	---	---	---	---	---	---	---	----

Running time?

When just one item is
out of place...

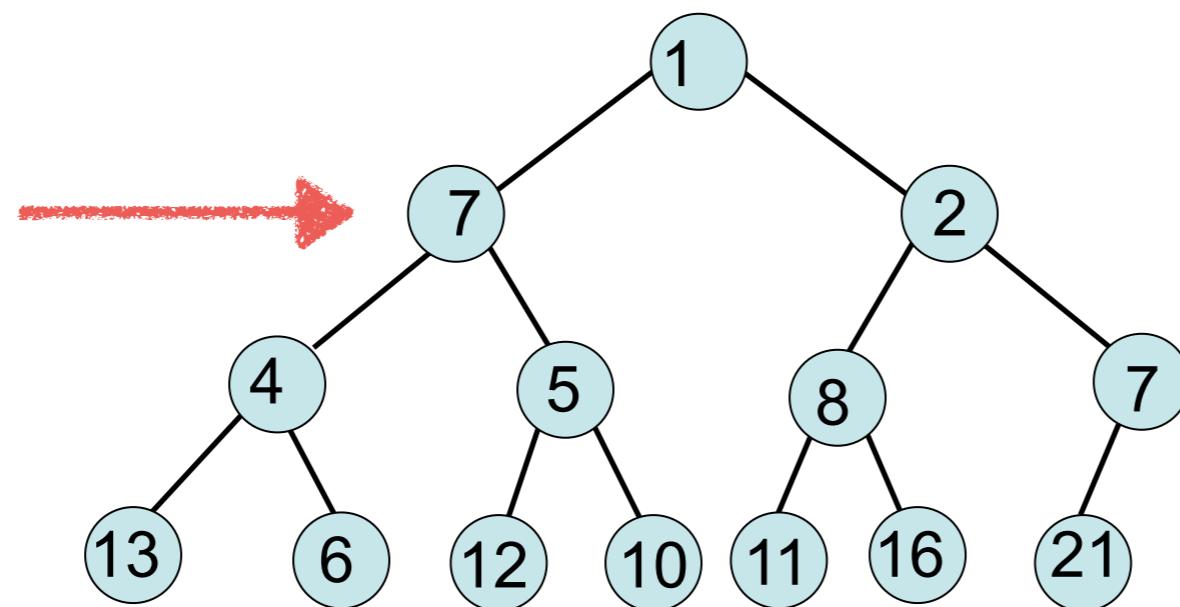
MIN-HEAPIFY(A , 2)

exchange $A[2]$ with $A[4]$

MIN-HEAPIFY(A , 4)

exchange $A[4]$ with $A[9]$

Assume trees rooted at $\text{LEFT}(i)$ and $\text{RIGHT}(i)$ are min-heaps



Maintaining the heap order property

MIN-HEAPIFY(A, i)

l = LEFT(i)

r = RIGHT(i)

if l <= A.heap-size **and** A[l] < A[i]

smallest = l

else smallest = i

if r <= A.heap-size **and** A[r] < A[smallest]

smallest = r

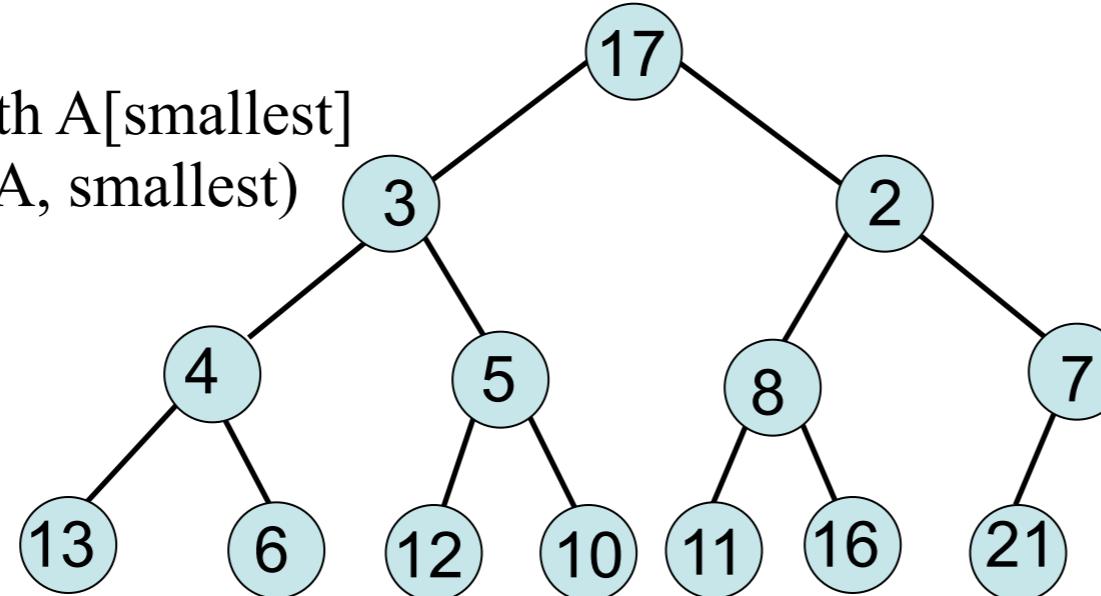
if smallest ≠ i

exchange A[i] with A[smallest]

MIN-HEAPIFY(A, smallest)

MIN-HEAPIFY(A,7)

percolateDown



1	2	3	4	5	6	7	8	9	10	11	12	13	14
17	3	2	4	5	8	7	13	6	12	10	11	16	21
2	3	17	4	5	8	7	13	6	12	10	11	16	21
2	3	7	4	5	8	17	13	6	12	10	11	16	21

44

Back to being a min binary heap!

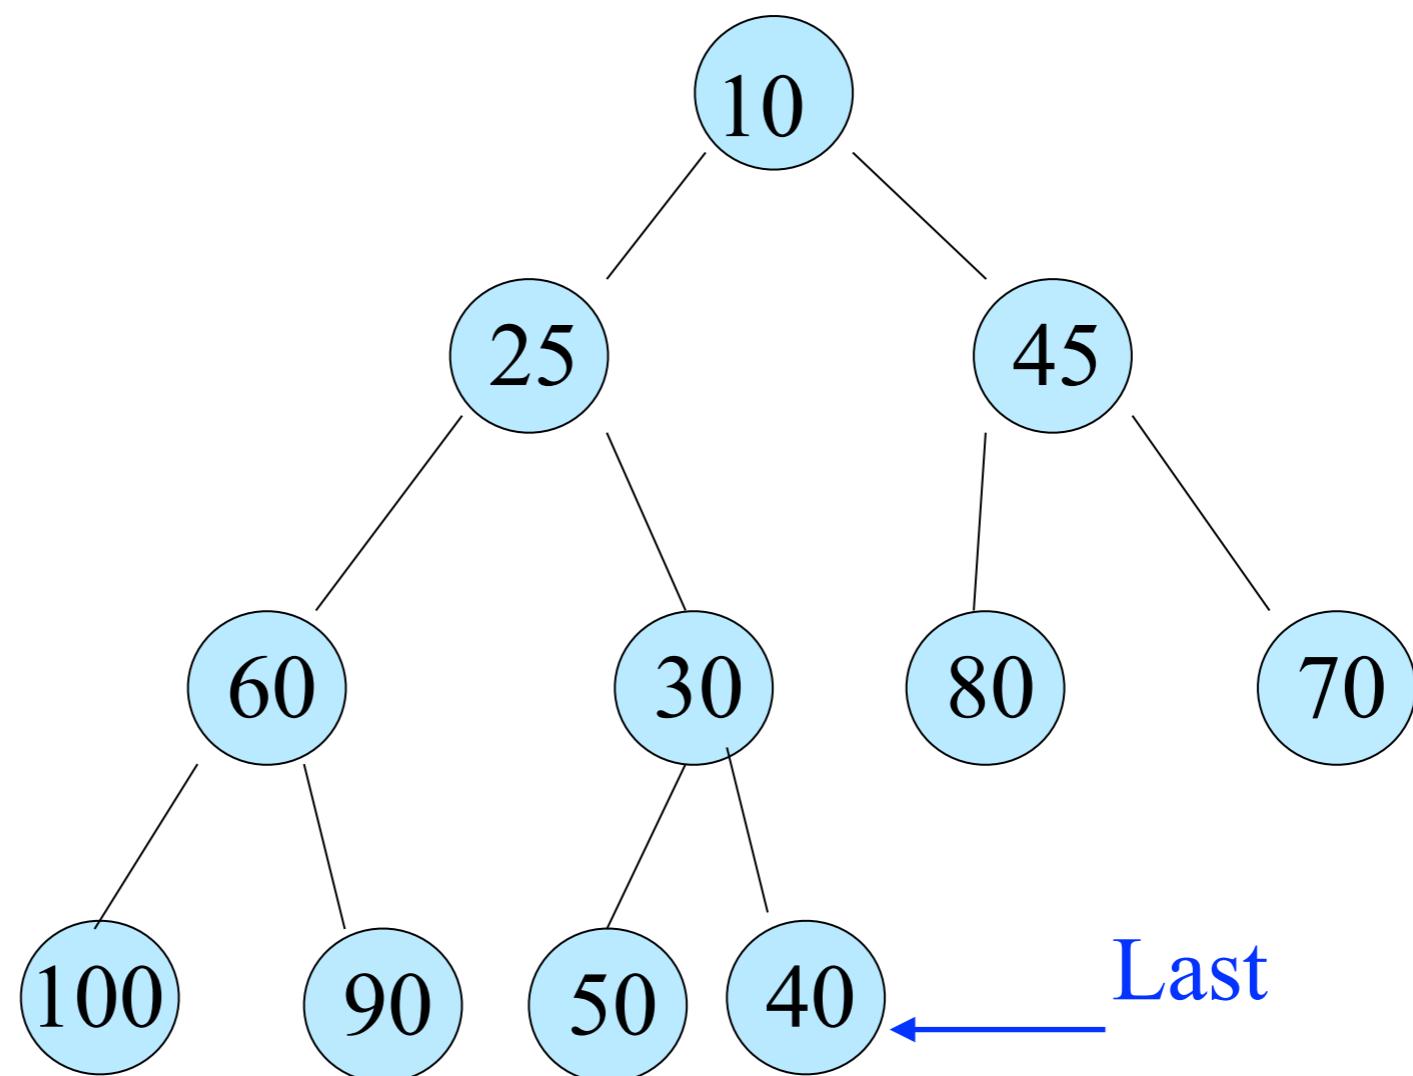
Heap-Extract-Min

HEAP-EXTRACT-MIN(A)

```
if A.heap-size < 1  
    error "heap underflow"  
min = A[1]  
A[1] = A[A.heap-size]  
A.heap-size = A.heap-size - 1  
MIN-HEAPIFY(A, 1)  
return min
```

Place last item in hole and percolate down

Delete Min



10	25	45	60	30	80	70	100	90	50	40	
1	2	3	4	5	6	7	8	9	46	10	11

Percolate down by swapping with smaller child.

MIN-HEAPIFY(A, i)

 l = LEFT(i)

 r = RIGHT(i)

if l <= A.heap-size **and** A[l] < A[i]

 smallest = l

else smallest = i

if r <= A.heap-size **and** A[r] < A[smallest]

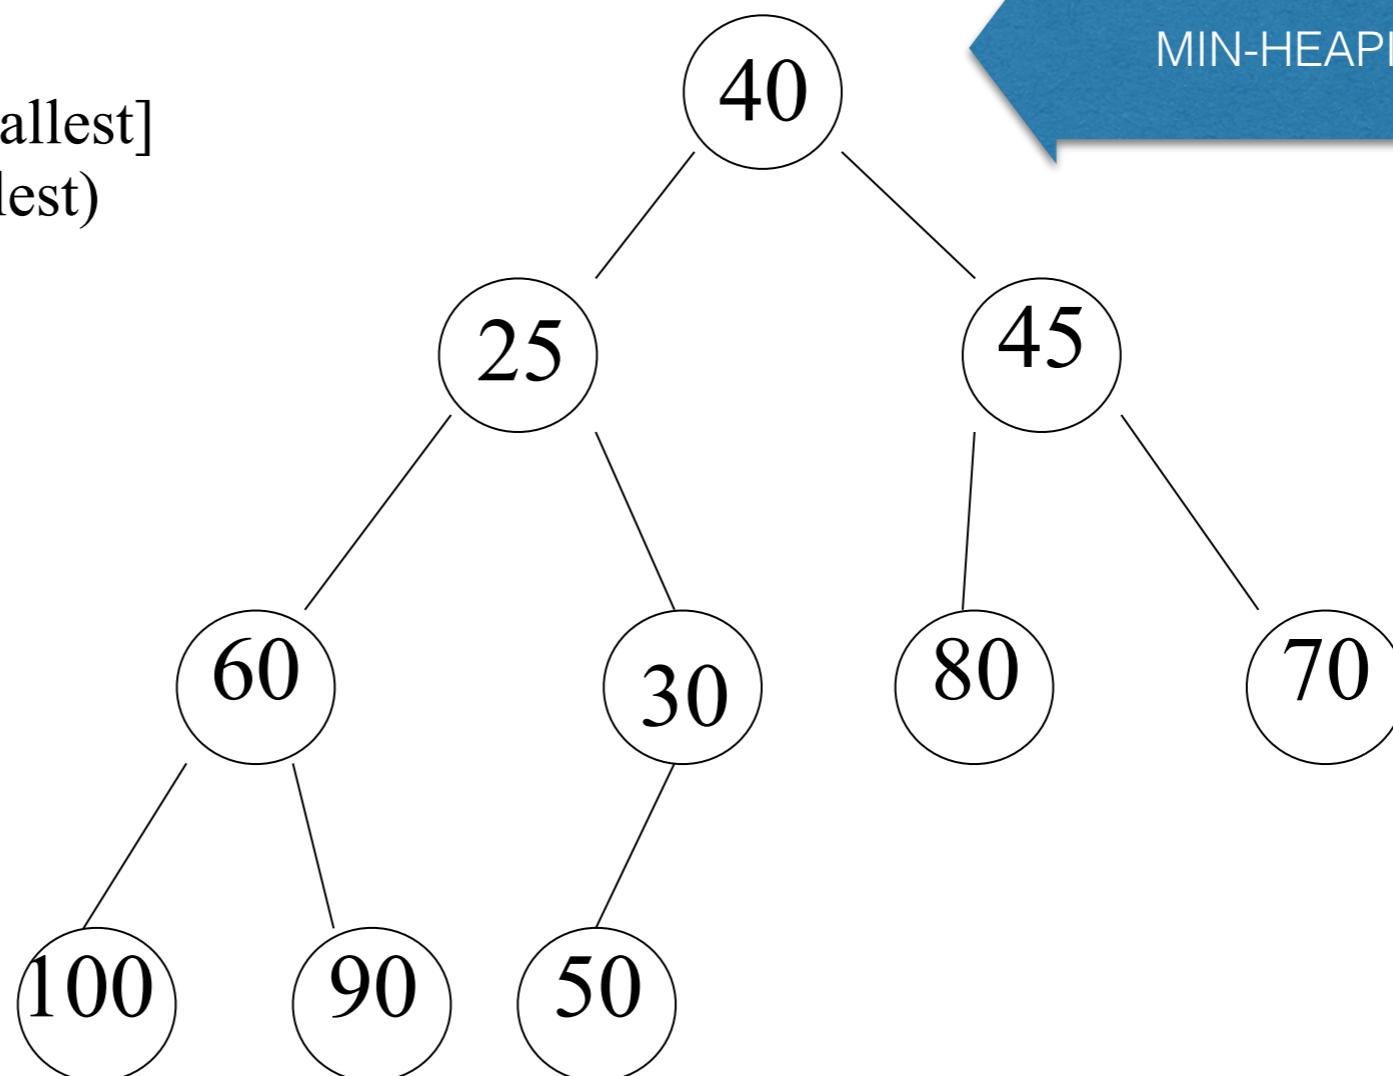
 smallest = r

if smallest ≠ i

 exchange A[i] with A[smallest]

 MIN-HEAPIFY(A, smallest)

MIN-HEAPIFY(A, 1)



40 25 45 60 30 80 70 100 90 50
1 2 3 4 5 6 7 8 9 10

Percolate down by swapping with smaller child.

MIN-HEAPIFY(A, i)

l = LEFT(i)

r = RIGHT(i)

if l <= A.heap-size **and** A[l] < A[i]

smallest = l

else smallest = i

if r <= A.heap-size **and** A[r] < A[smallest]

smallest = r

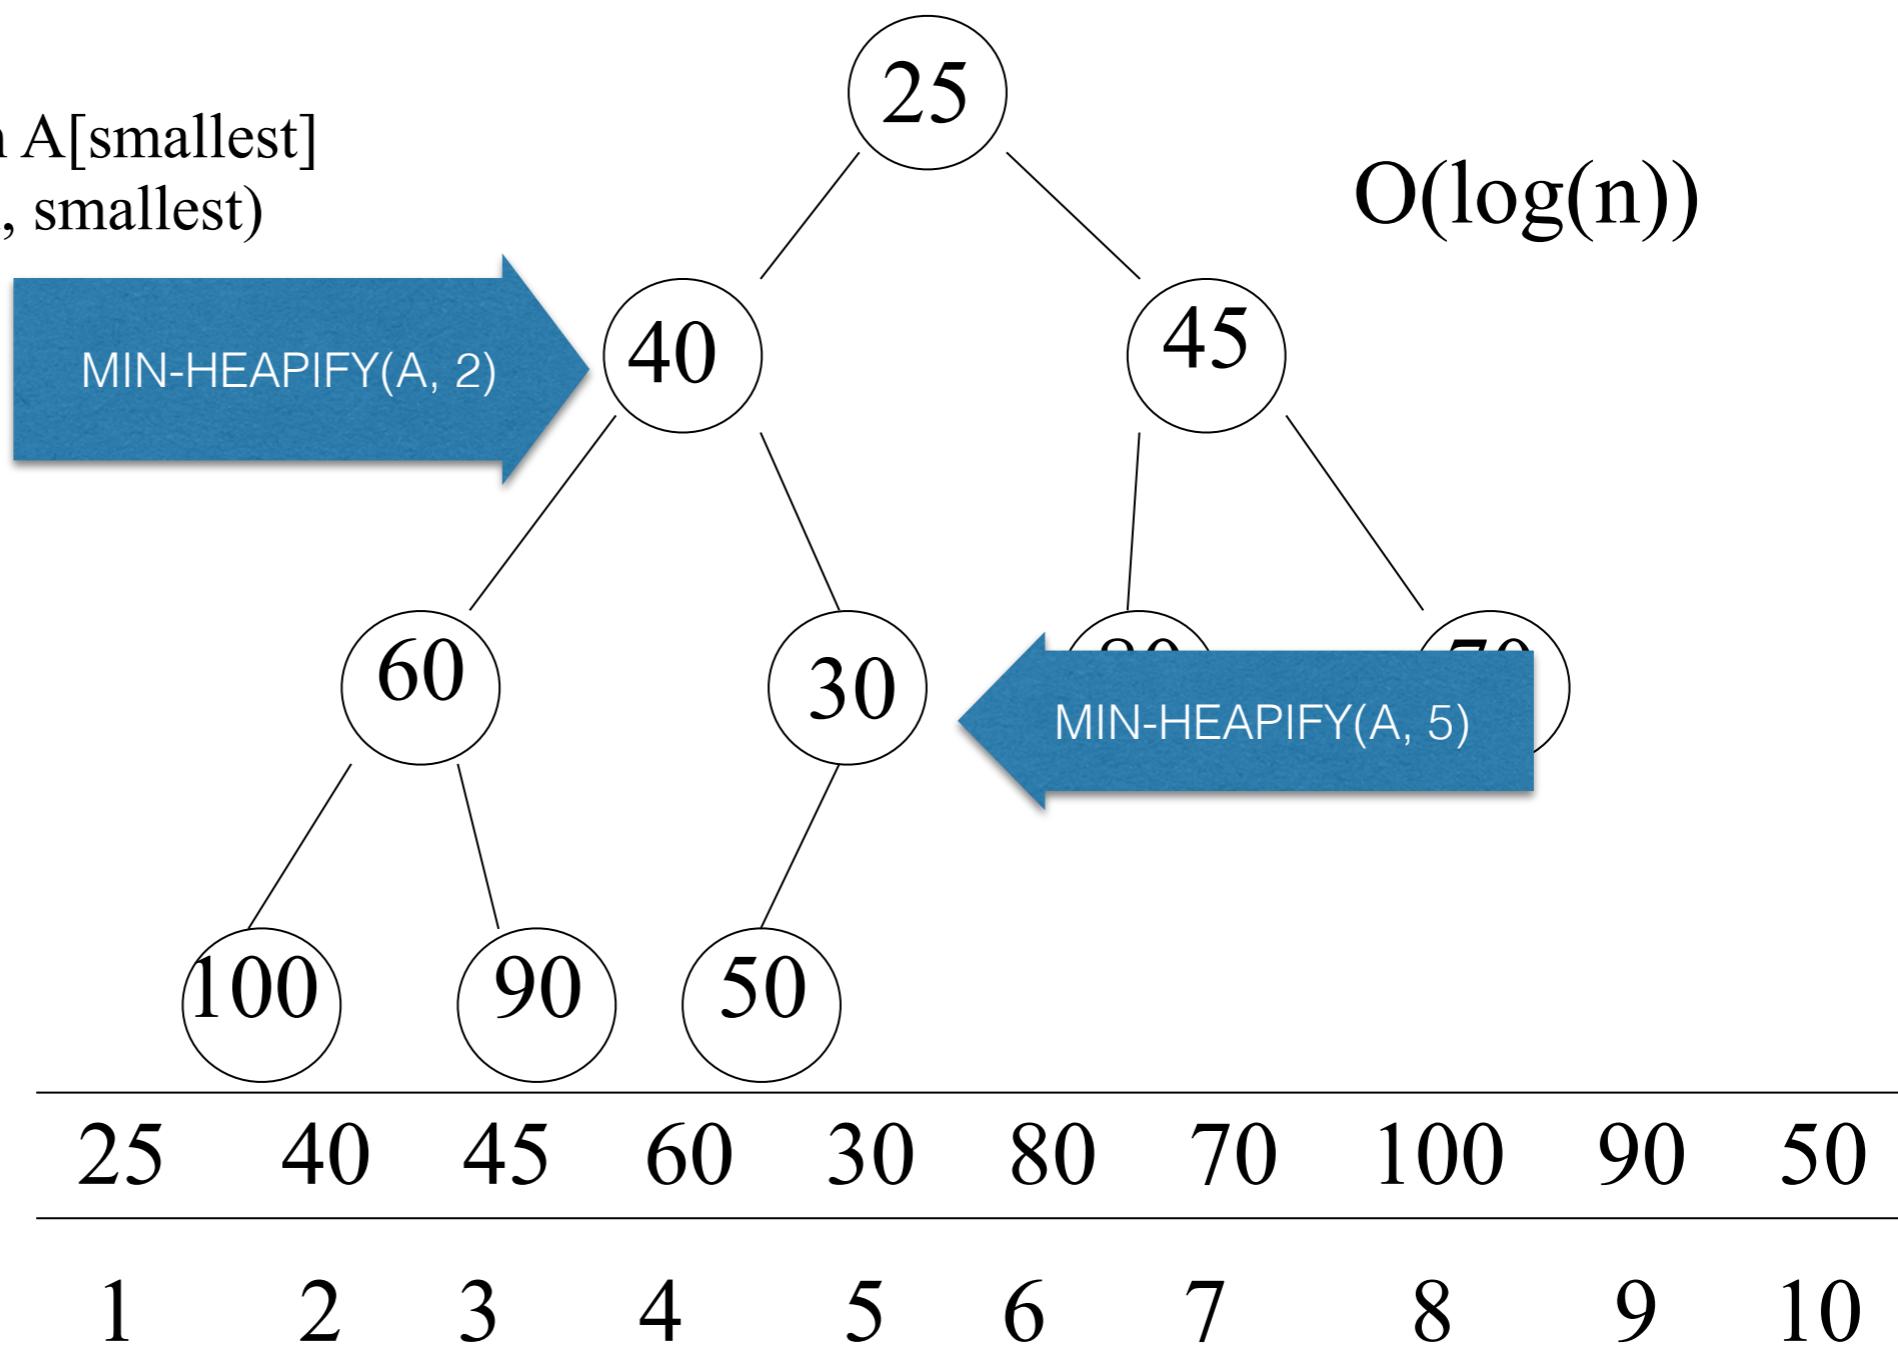
if smallest ≠ i

exchange A[i] with A[smallest]

MIN-HEAPIFY(A, smallest)

Time?

O(log(n))



Heap-Decrease-Key

HEAP-DECREASE-KEY(A, i, key)

if key > A[i]

error “new key is larger”

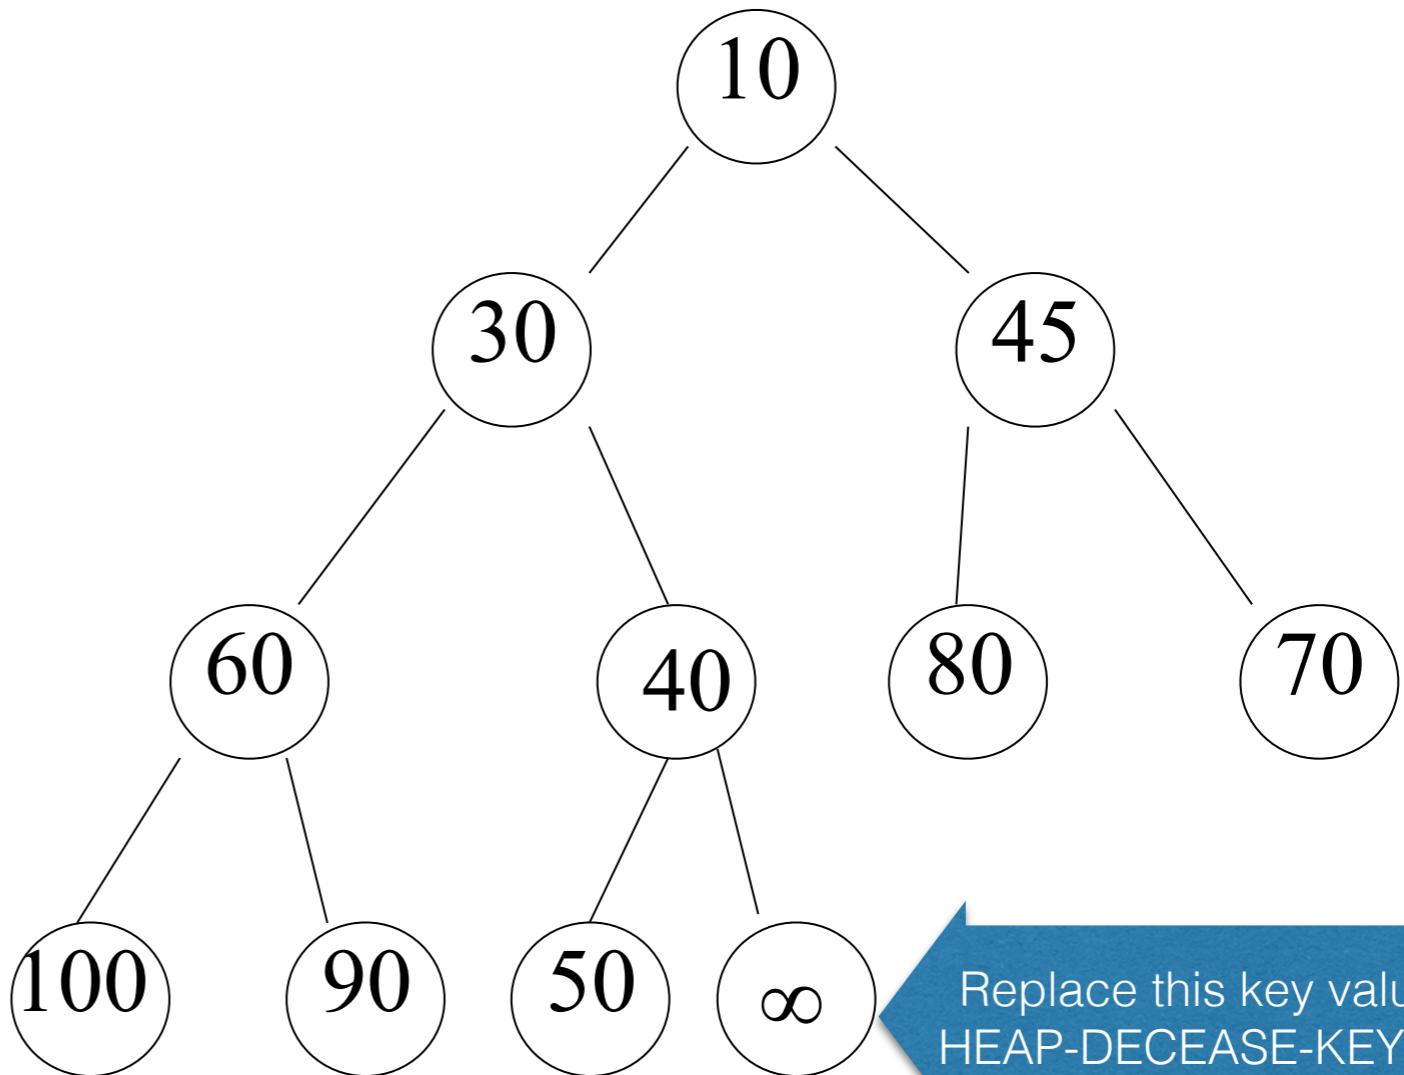
A[i] = key

while i > 1 and A[PARENT(i)] > A[i]

exchange A[i] with A[PARENT(i)]

i = PARENT(i)

Percolate up



10	30	45	60	40	80	70	100	90	50	infinity
1	2	3	4	5	6	7	8	9	10	11

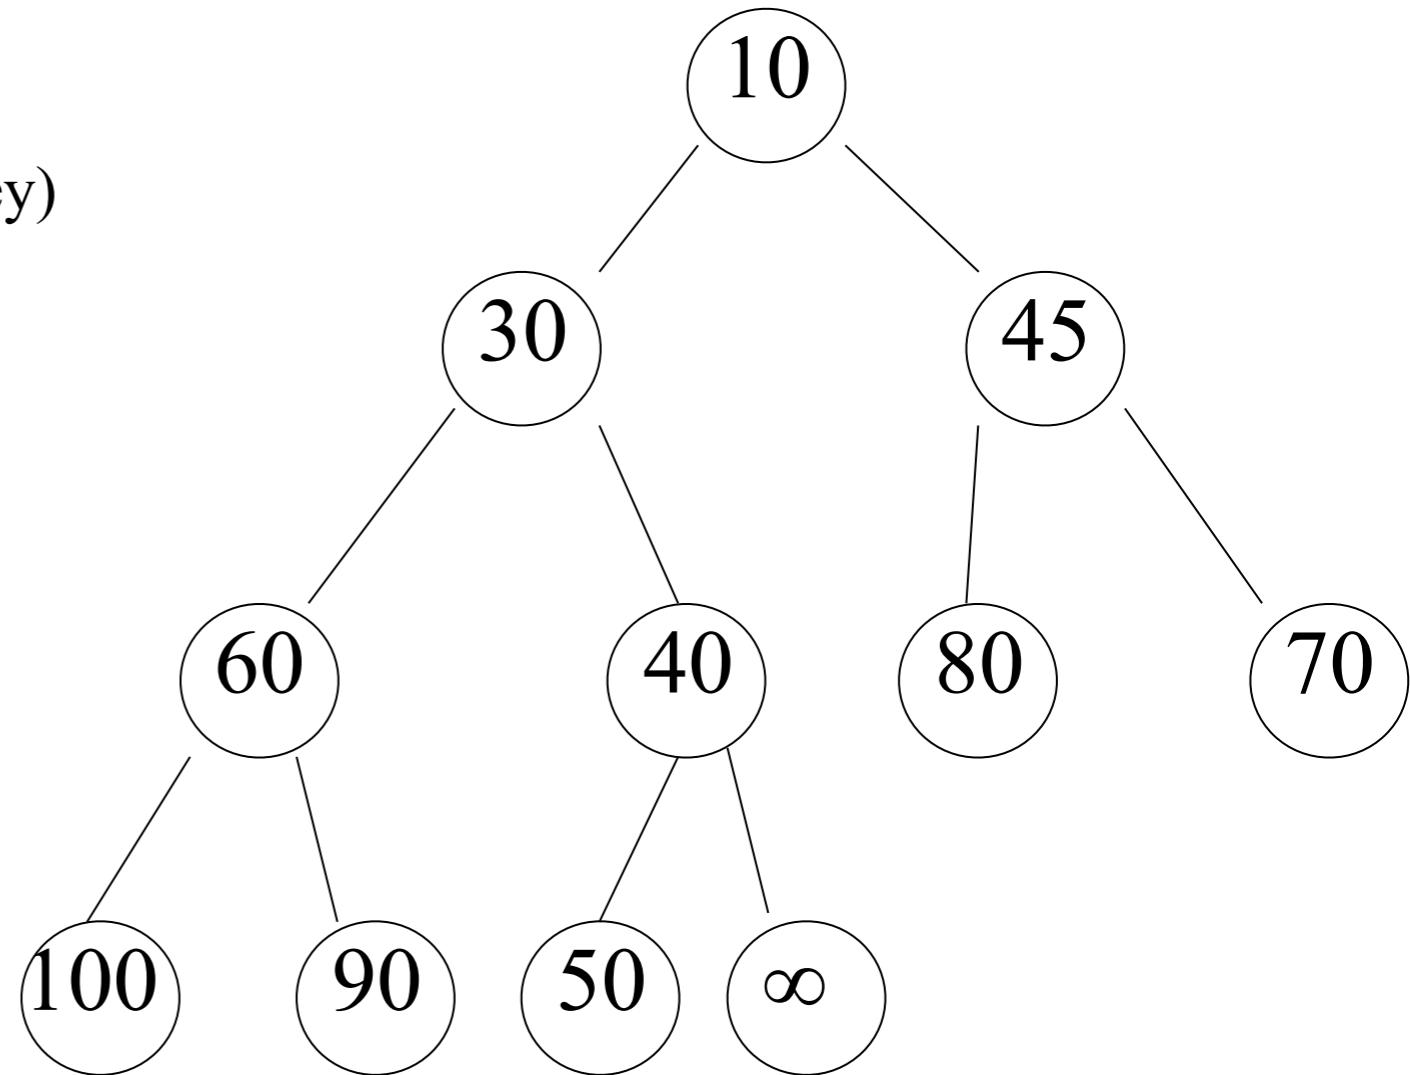
Min-Heap-Insert

MIN-HEAP-INSERT(A, key)

A.heap-size = A.heap-size + 1

A[A.heap-size] = ∞

HEAP-DECREASE-KEY(A, A.heap-size, key)



10	30	45	60	40	80	70	100	90	50	∞
1	2	3	4	5	6	7	8	9	10	11

HEAP-DECREASE-KEY(A, i, key)

Percolate up

HEAP-DECREASE-KEY(A, 11, 25)

if key > A[i]

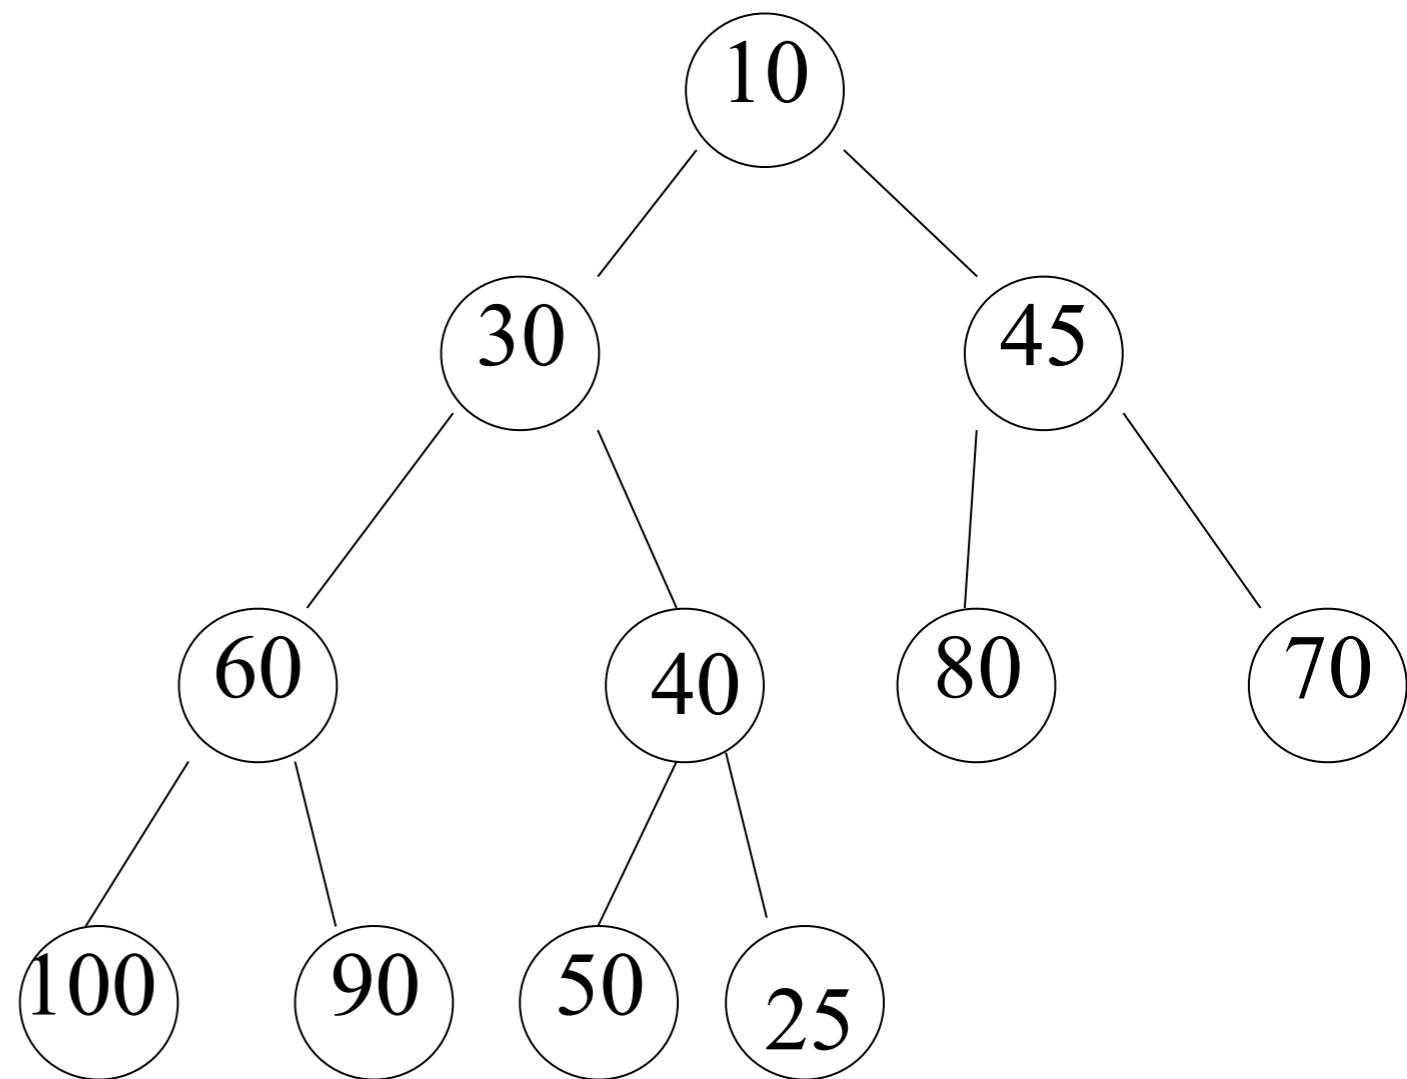
error “new key is larger”

A[i] = key

while i > 1 **and** A[PARENT(i)] > A[i]

exchange A[i] with A[PARENT(i)]

 i = PARENT(i)



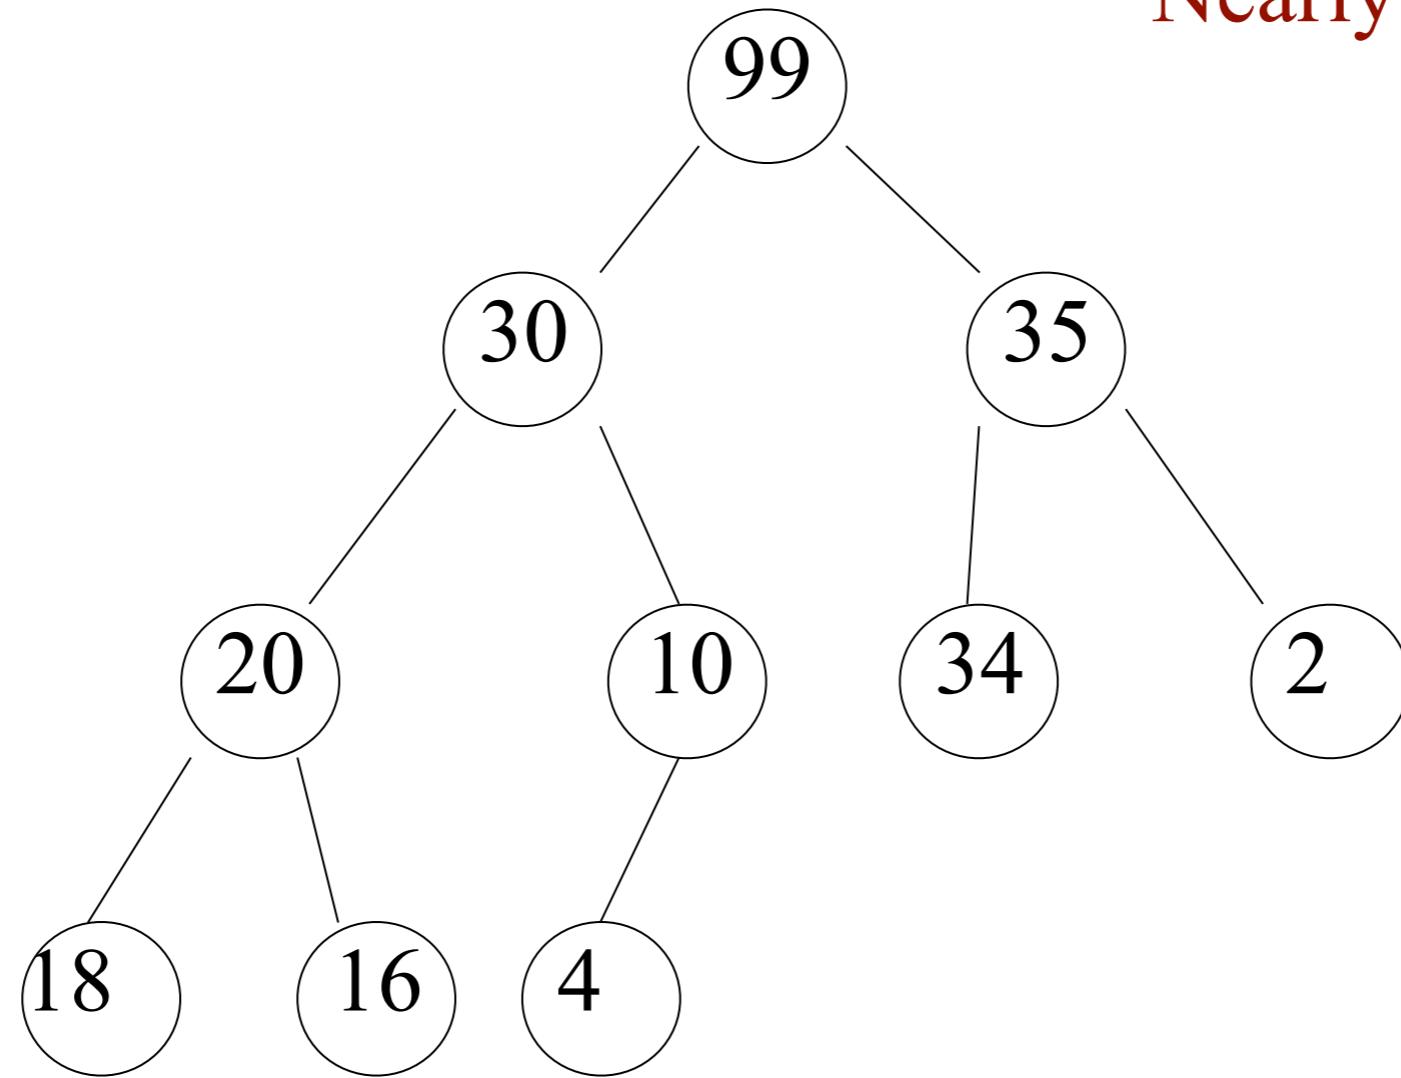
10	30	45	60	40	80	70	100	90	50	25
1	2	3	4	5	6	7	8	9	10	11

Time for insertion? O(log(n)) Random input average 2.607 ⁵³ comparisons!

A Max Heap Example

Max Heap

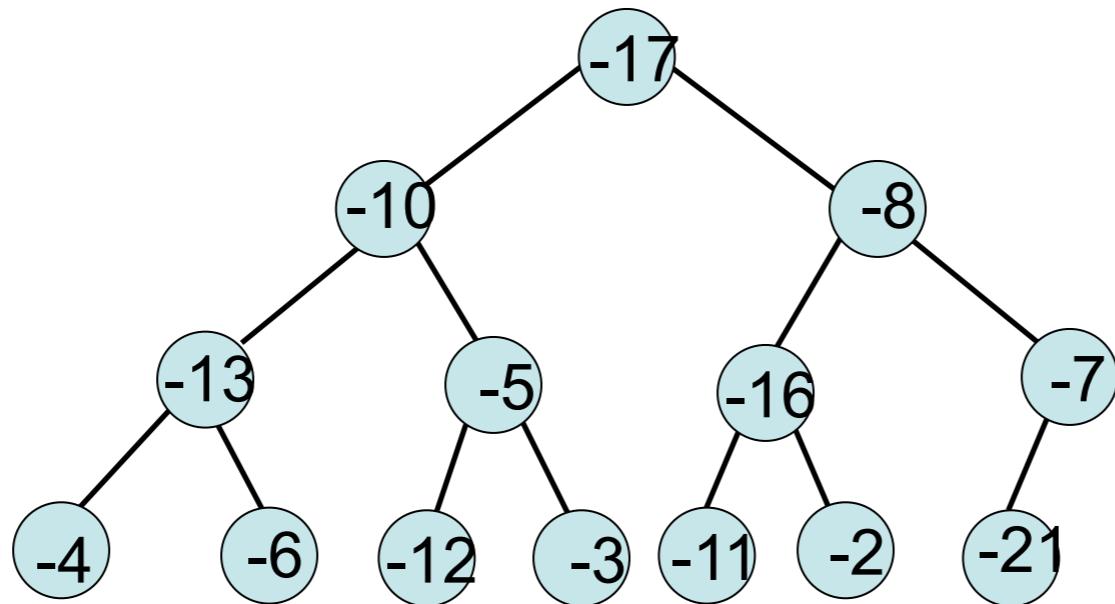
Nearly Complete Tree



How can we turn A into a max heap?

	1	2	3	4	5	6	7	8	9	10	11	12	13	14
A	-17	-10	-8	-13	-5	-16	-7	-4	-6	-12	-3	-11	-2	-21

How can we turn A into a max heap?



Where to start?
How long does it take?

	1	2	3	4	5	6	7	8	9	10	11	12	13	14
A	-17	-10	-8	-13	-5	-16	-7	-4	-6	-12	-3	-11	-2	-21

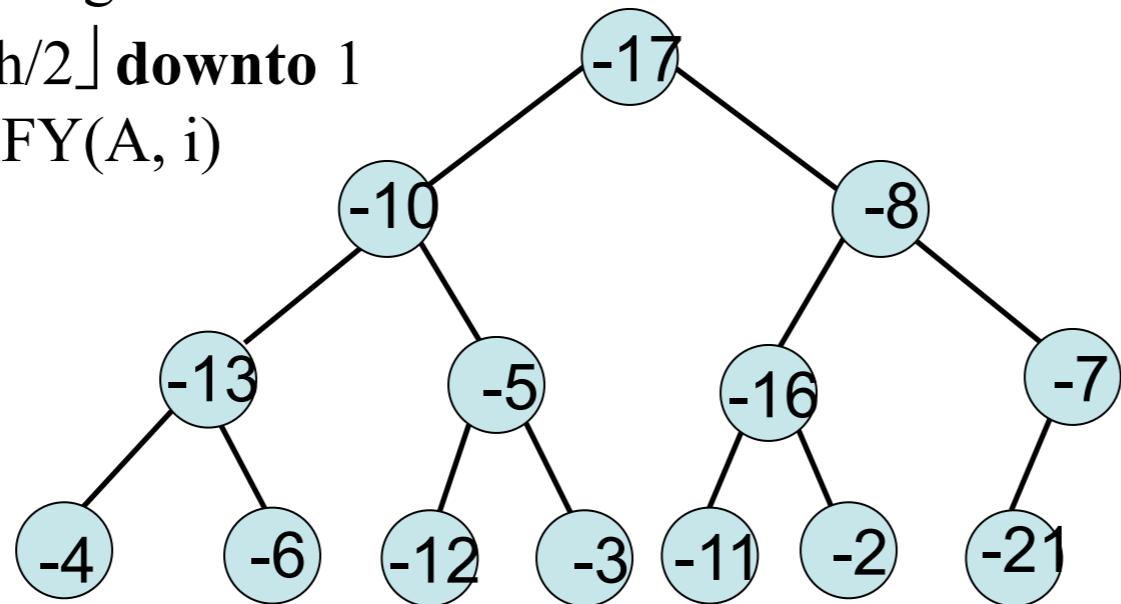
Building a max-heap

BUILD-MAX-HEAP(A)

$A.\text{heap-size} = A.\text{length}$

for $i = \lfloor A.\text{length}/2 \rfloor$ **downto** 1

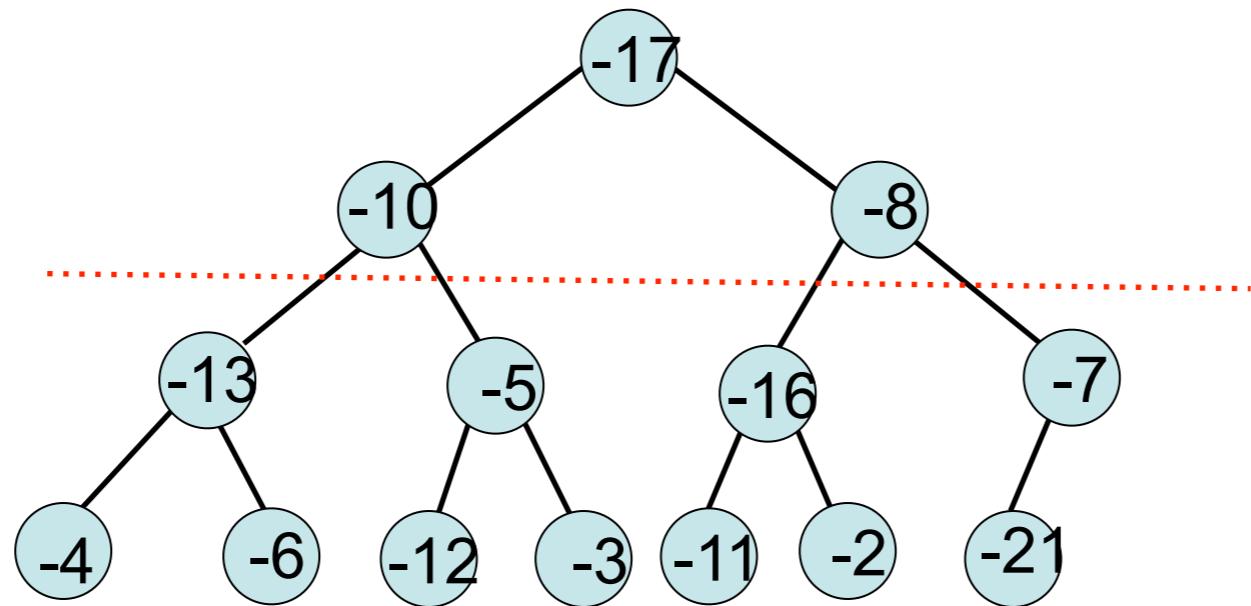
MAX-HEAPIFY(A, i)



1	2	3	4	5	6	7	8	9	10	11	12	13	14
-17	-10	-8	-13	-5	-16	-7	-4	-6	-12	-3	-11	-2	-21

BUILD-MAX-HEAP(A)

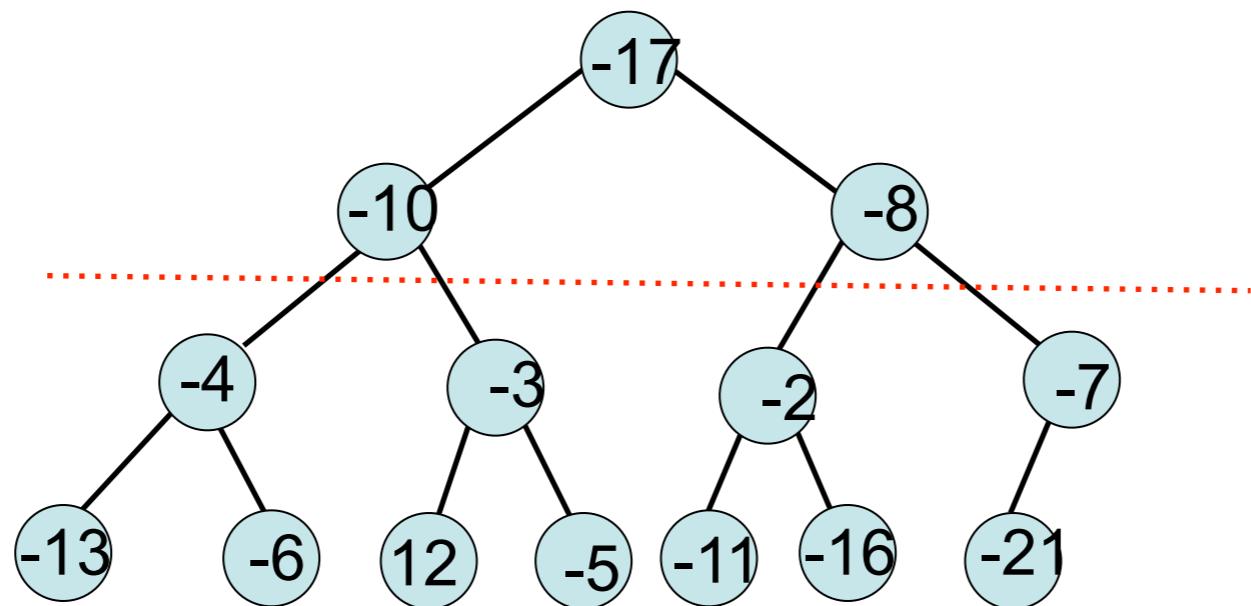
```
A.heap-size = A.length  
for i = ⌊A.length/2⌋ downto 1  
    MAX-HEAPIFY(A, i)
```



percolateDown all trees
below the dotted line

BUILD-MAX-HEAP(A)

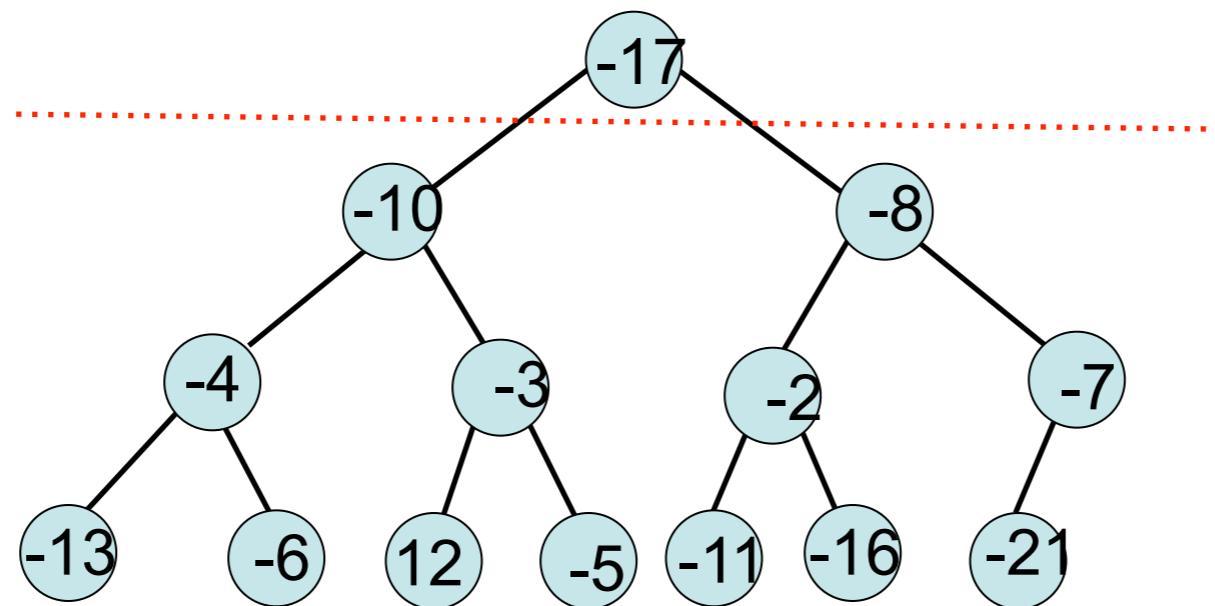
```
A.heap-size = A.length  
for i = ⌊A.length/2⌋ downto 1  
    MAX-HEAPIFY(A, i)
```



Done percolating
down all trees below the
dotted line

BUILD-MAX-HEAP(A)

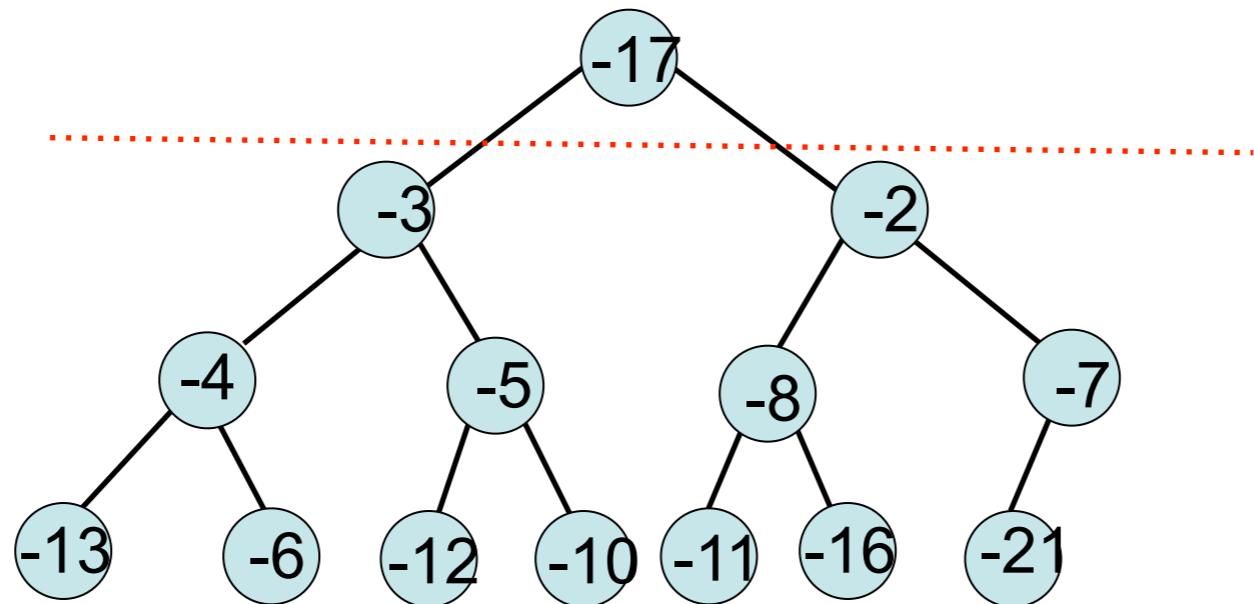
```
A.heap-size = A.length  
for i = ⌊A.length/2⌋ downto 1  
    MAX-HEAPIFY(A, i)
```



percolating down **all**
trees below the dotted line

BUILD-MAX-HEAP(A)

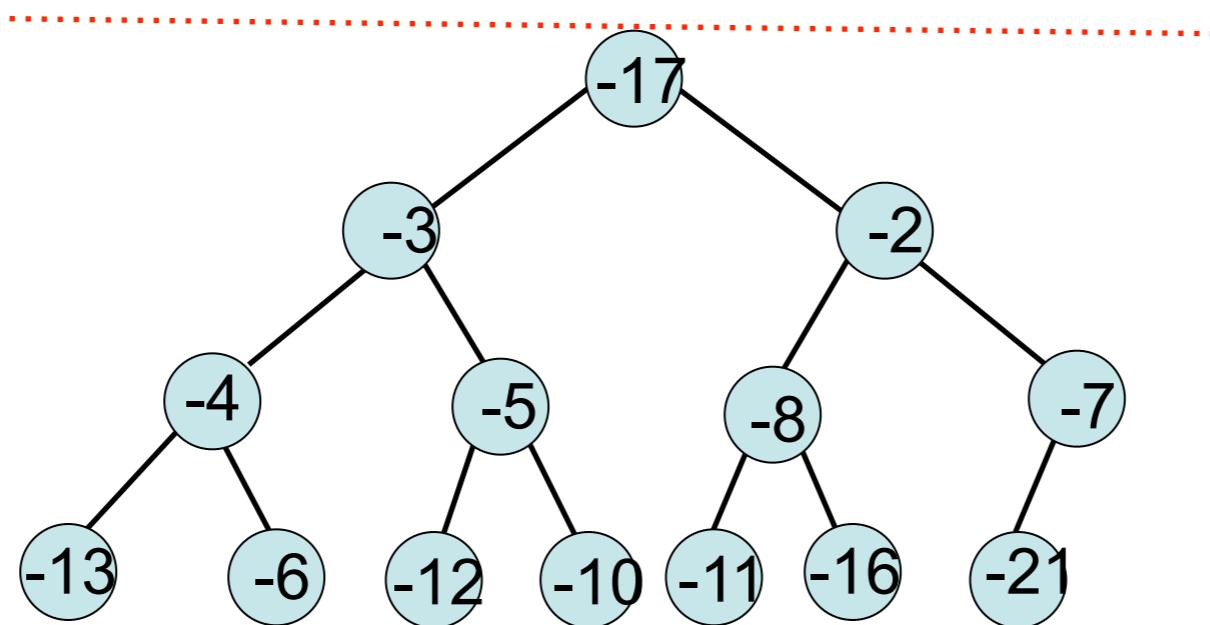
```
A.heap-size = A.length  
for i = ⌊A.length/2⌋ downto 1  
    MAX-HEAPIFY(A, i)
```



Done percolating
down all trees below the
dotted line

BUILD-MAX-HEAP(A)

```
A.heap-size = A.length  
for i = ⌊A.length/2⌋ downto 1  
    MAX-HEAPIFY(A, i)
```

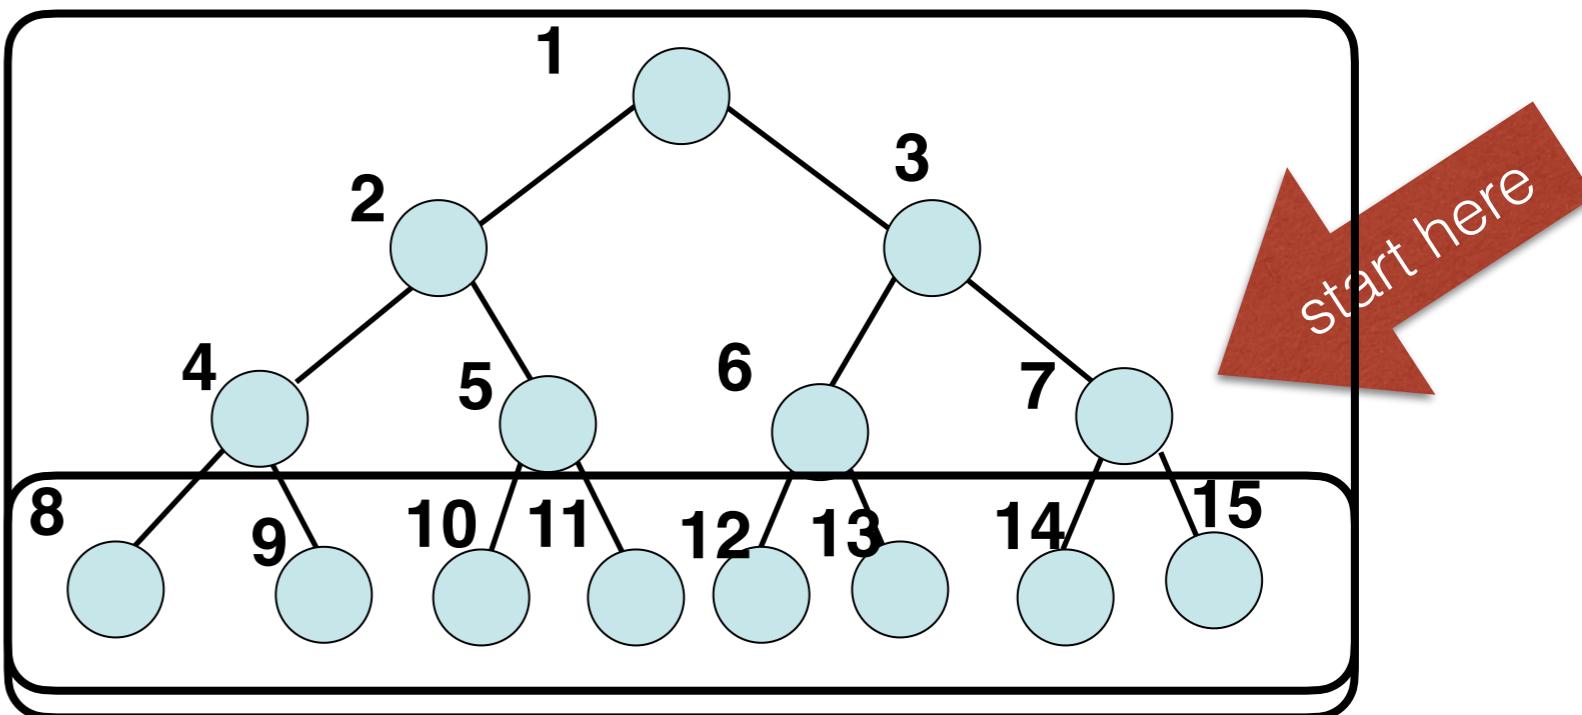


percolateDown all trees
below the dotted line

Can we prove
BUILD-MAX-HEAP works?

BUILD-MAX-HEAP(A)

```
A.heap-size = A.length  
for i = ⌊A.length/2⌋ downto 1  
    MAX-HEAPIFY(A, i)
```



Loop Invariant: At the **start** of each iteration of the for loop, each node $i + 1, i + 2, \dots, n$ is the root of a max-heap.

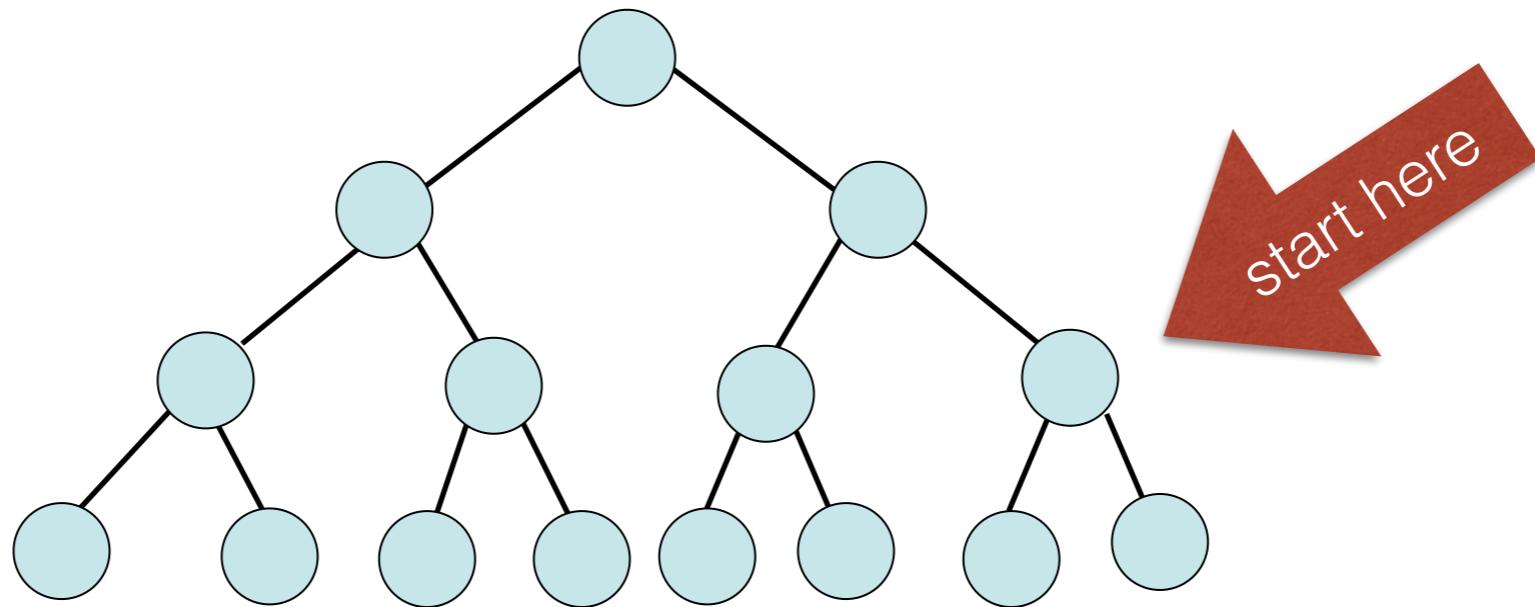
Initialization: Each node $\lfloor n/2 \rfloor + 1, \lfloor n/2 \rfloor + 2, \dots, n$ is a leaf, which is the root of a max-heap. Since $i = \lfloor n/2 \rfloor$ before the first iteration of the for loop, the invariant is initially true.

Maintenance: Children of node i are indexed higher than i , so by the loop invariant, they are both max-heaps. Correctly assuming that $i+1, i+2, \dots, n$ are all roots of max-heaps, MAX-HEAPIFY makes node i a max-heap root.

Decrementing i re-establishes the loop invariant at each iteration.

Termination: When the loop terminates, $i = 0$. By the loop invariant, each node, notably node 1, is the root of a max-heap.

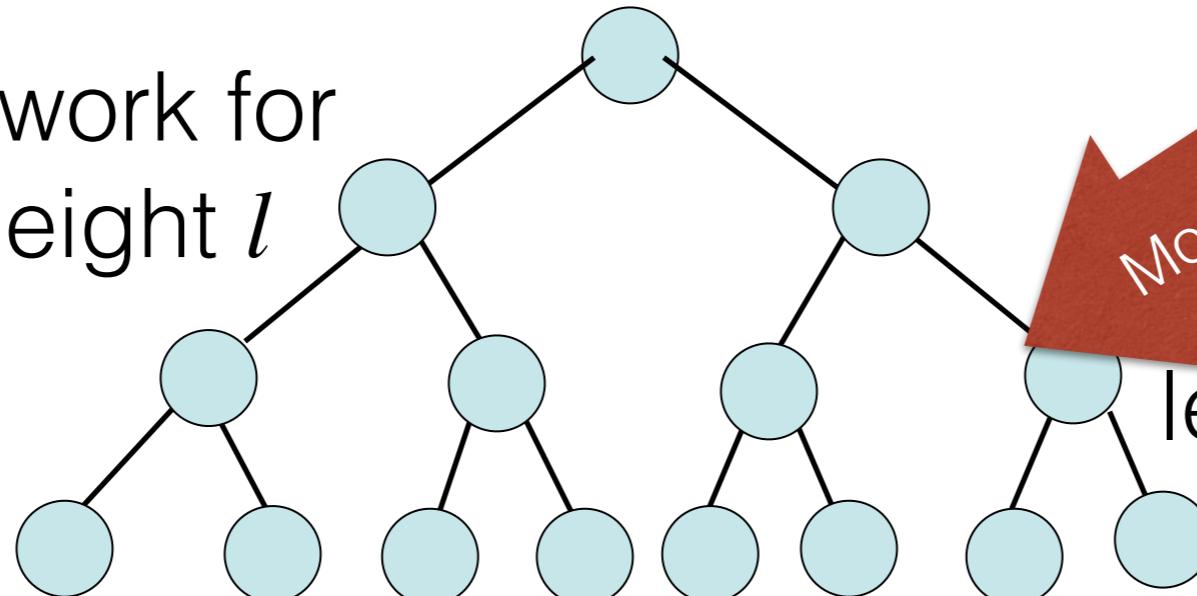
How long does BUILD-MAX-HEAP take?



Does it really take that long?

- Each percolate down takes $O(\log n)$ time. We call percolate down $\sim n/2$ times. Therefore we know BUILD-MAX-HEAP(A) takes $O(n \log n)$ time.
- Our typical analysis is that if we do x operations that take time at most y , then the running time is $O(x \cdot y)$
- Did we overestimate too much?

$O(l)$ amount of work for any node at height l



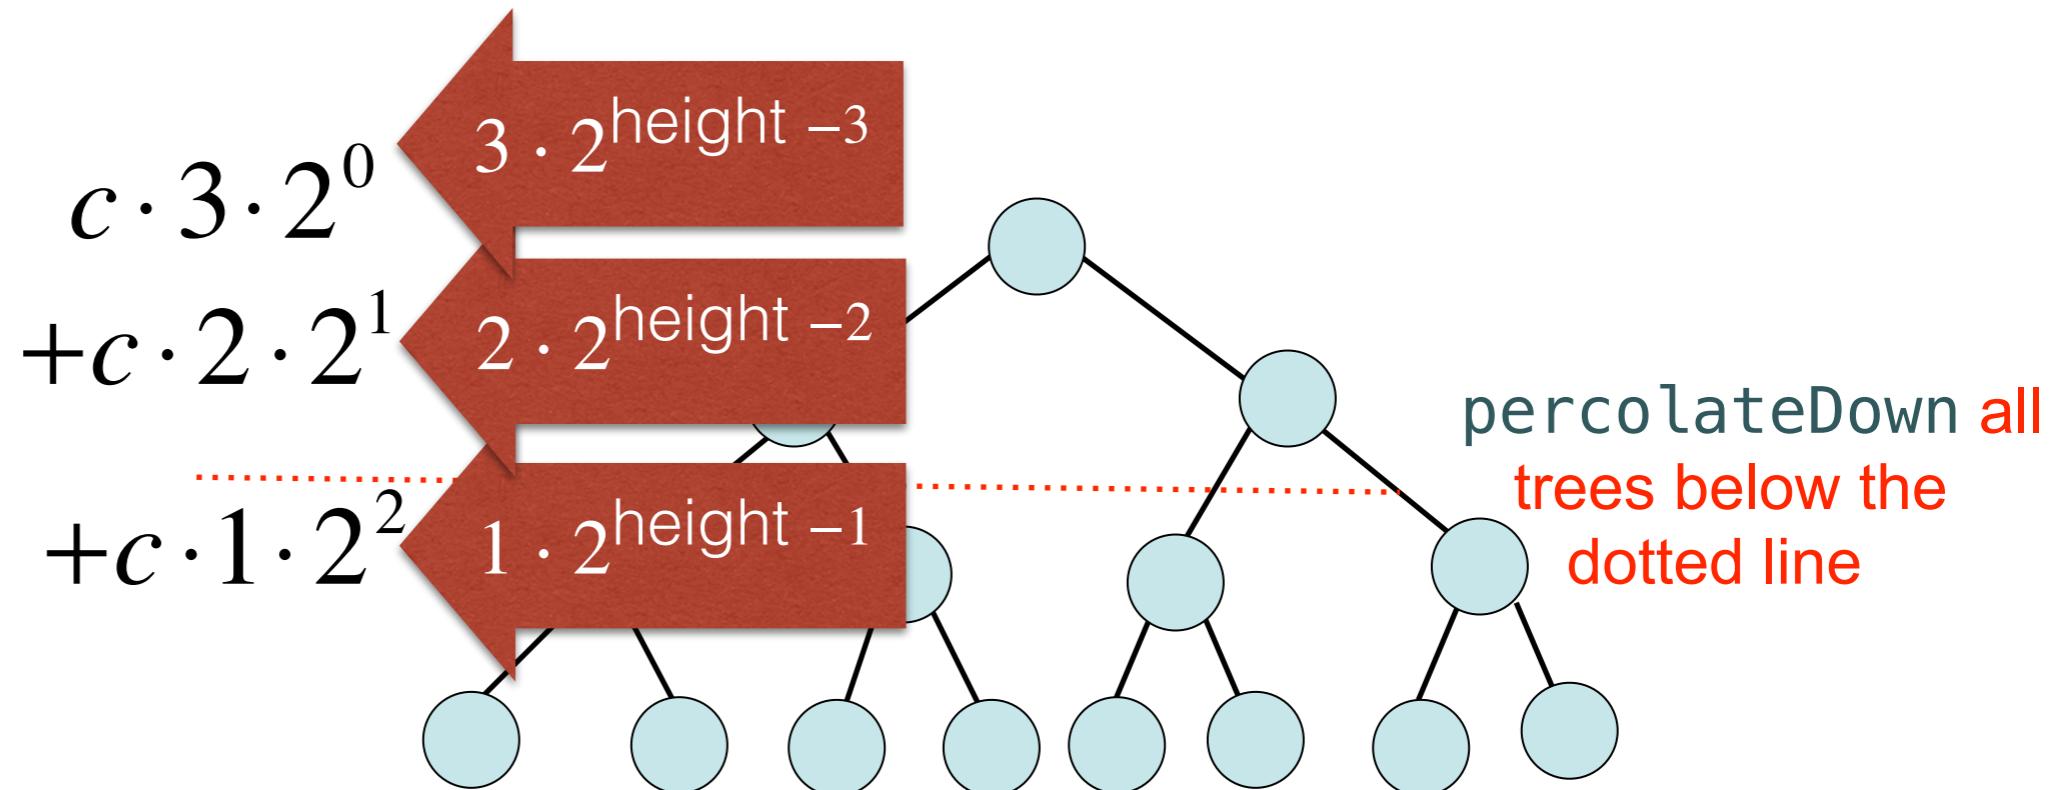
Most nodes have small height
each node on this level takes a constant amount of work

“Amortized analysis: Determine worst-case running time of a sequence of data structure operations as a function of the input size.”

“Aggregate method: Sum up sequence of operations, weighted by their cost.”

How many operations?

<i>height</i>	<i>depth</i>	max # nodes
3	0	$1 = 2^0$
2	1	$2 = 2^1$
1	2	$4 = 2^2$
0	3	$7 \leq 2^3$



$$T(n) \leq c \cdot 1 \cdot 2^2 + c \cdot 2 \cdot 2^1 + c \cdot 3 \cdot 2^0$$

$$T(n) \leq c \sum_{i=1}^h 2^{h-i} \cdot i$$

i = node height

h = height of tree

h = height of tree

node height
depth max # nodes

h	0	$1 = 2^0$
$h-1$	1	$2 = 2^1$
$h-2$	2	$4 = 2^2$
$h-3$	3	$8 = 2^3$
$h-i$	i	2^i
1	$h-1$	2^{h-1}

How many operations?

$$c \cdot h \cdot 2^0$$

$$c \cdot (h-1) \cdot 2^1$$

$$c \cdot (h-2) \cdot 2^2$$

$$c \cdot (h-3) \cdot 2^3$$

$$c \cdot (h-i) \cdot 2^i$$

$$c \cdot (1) \cdot 2^{h-1}$$

$i = \text{node height}$

$$T(n) \leq c \sum_{i=1}^h 2^{h-i} \cdot i \leq c 2^{\log n} \sum_{i=1}^h 2^{-i} \cdot i \leq cn \sum_{i=1}^{\infty} 2^{-i} \cdot i$$

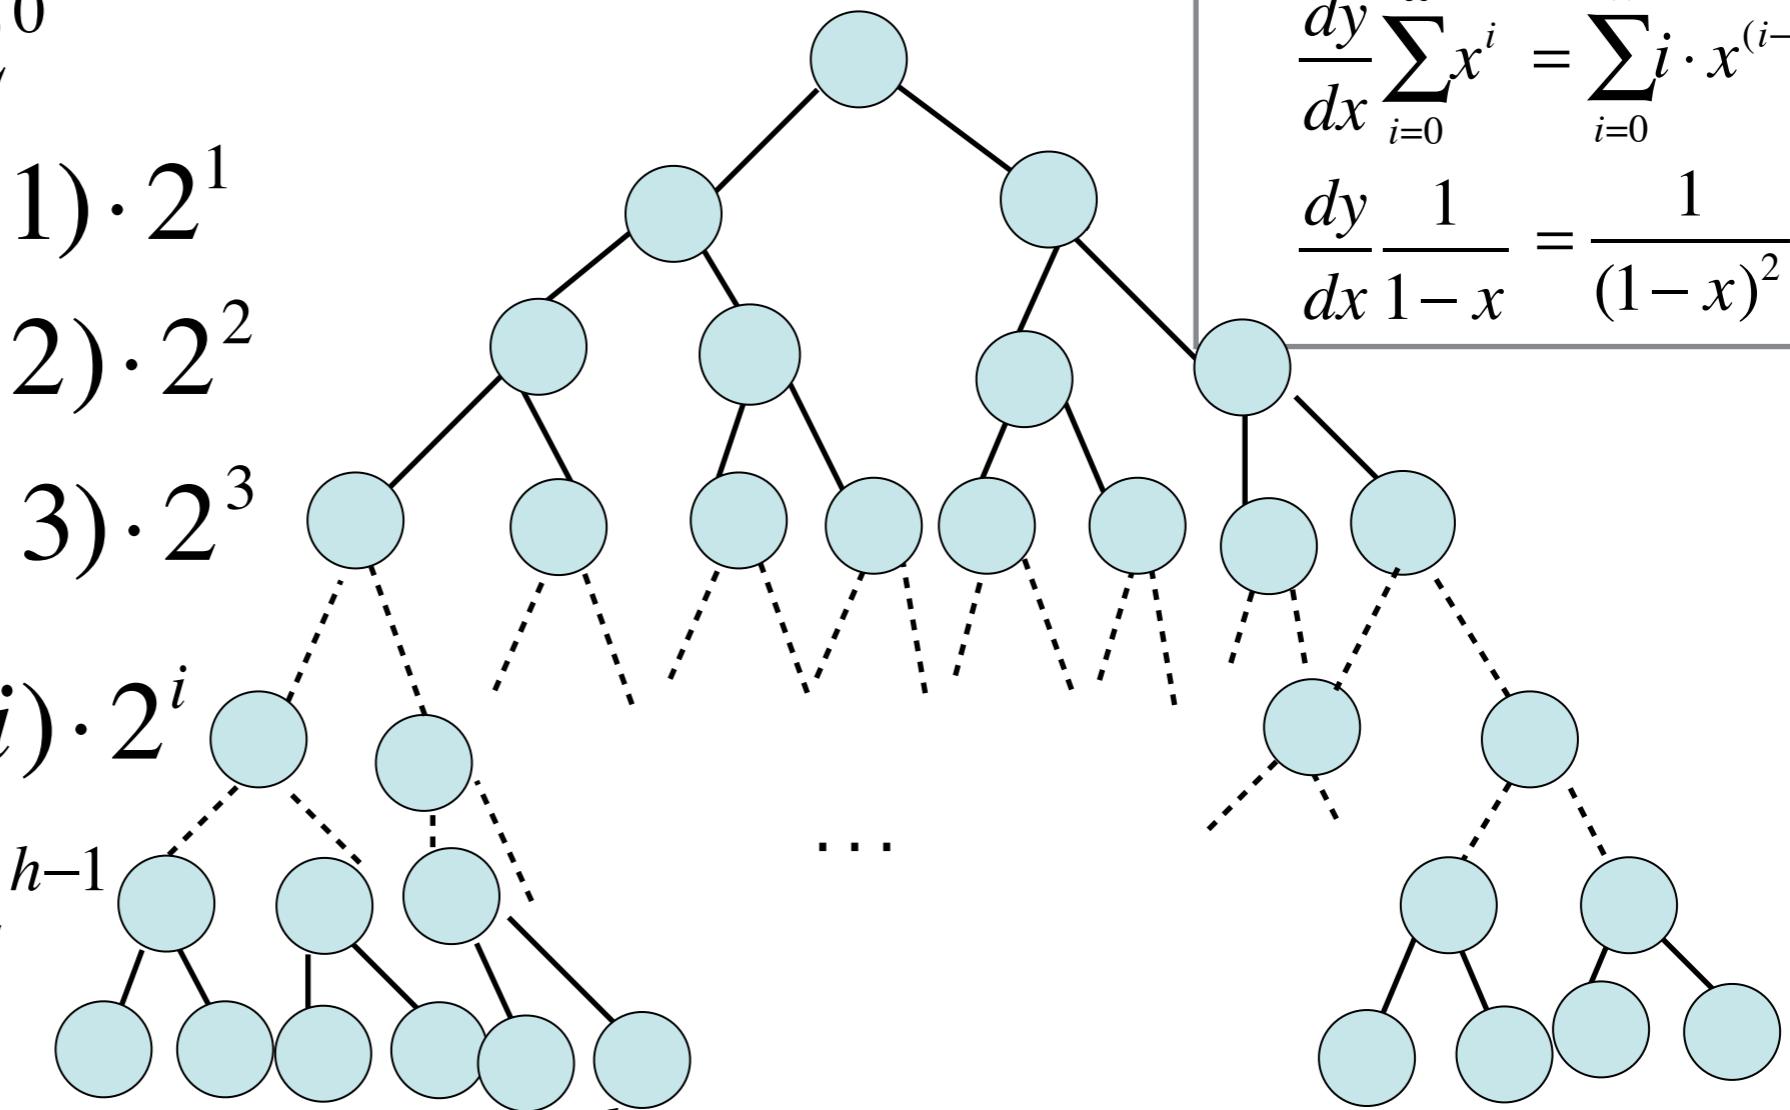
$$= cn \sum_{i=1}^{\infty} \left(\frac{1}{2}\right)^i \cdot i = \frac{cn}{2} \sum_{i=1}^{\infty} i \cdot \left(\frac{1}{2}\right)^{i-1} = \frac{cn}{2} \frac{1}{(1-1/2)^2} = O(n)$$

$|x| < 1 \rightarrow \text{workspace}$

$$\text{let } x = \frac{1}{2} \quad \sum_{i=0}^{\infty} x^i = \frac{1}{1-x}$$

$$\frac{dy}{dx} \sum_{i=0}^{\infty} x^i = \sum_{i=0}^{\infty} i \cdot x^{(i-1)}$$

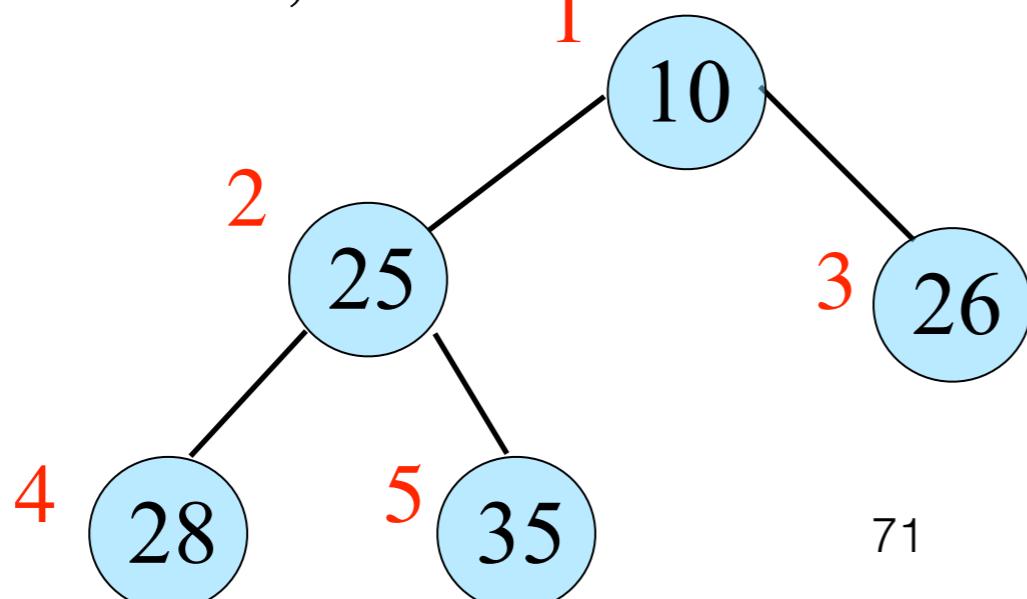
$$\frac{dy}{dx} \frac{1}{1-x} = \frac{1}{(1-x)^2}$$



Heap Sort

Given an input array, the heapsort algorithm acts as follows:

- Builds a max-heap from the array.
- Starting with the root (the maximum element), the algorithm places the maximum element into the correct place in the array by swapping it with the element in the last position in the array.
- “Discard” this last node (knowing that it is in its correct place) by decreasing the heap size, and calling MAX-HEAPIFY on the new (possibly incorrectly-placed) root.
- Repeat this “discarding” process until only one node (the smallest element) remains, and therefore is in the correct place in the array.



HEAPSORT(A)

HeapSort

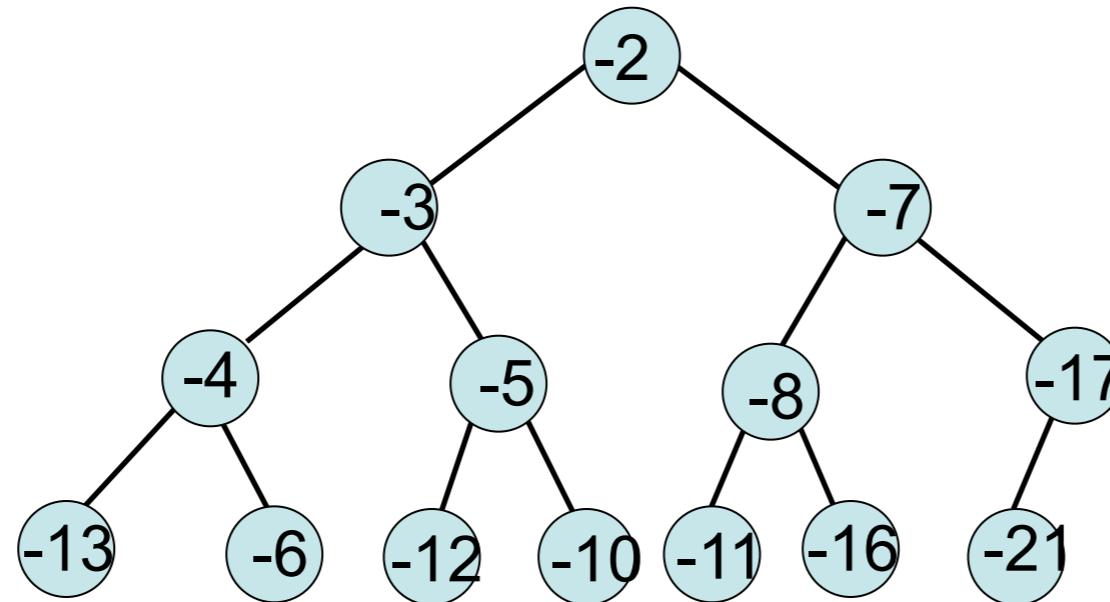
BUILD-MAX-HEAP(A)

for i = A.length **downto** 2

 exchange A[1] with A[i]

 A.heap-size = A.heap-size - 1

 MAX-HEAPIFY(A, 1)



And so on... till all the items are sorted

Heap-Sort

O(n log n) worst case time!

sort in place!